

1.	Prerequisites:	3
2.	Installation of Kafka – 90 Mins.....	8
3.	Installation Confluent Kafka (Local) – 30 Minutes	16
4.	Basic Kafka Operations - CLI (Topic) – 30 Mins	17
5.	Basic Kafka Operations - CLI (Producer & Consumer) – 30 Mins.....	21
6.	Zookeeper – 120 Minutes	23
7.	Kafka cluster – 90 Minutes.....	42
8.	Securing Kafka ACL – 60 Minutes	67
9.	Kafka Connector (File) - 60 Minutes.....	90
10.	Kafka Connector (JDBC) - 60 Minutes.....	99
12.	Errors	120
I.	{test=LEADER_NOT_AVAILABLE} (org.apache.kafka.clients.NetworkClient)	120
13.	LOG verification + segment Sizing	122
14.	Annexure Code:.....	123
II.	DumplogSegment.....	123
III.	Resources	124

Hardware:

8 GB RAM , 30 GB HDD , Centos 7 or above OS. Access to internet.

Software Inventory:

2 Kafka – Administration

- Zookeeper Version: apache-zookeeper-3.8.0-bin.tar
- Apache kafka : 2.13-3.2.1
- JDK 11.0.16
- Eclipse for Linux. (Any Latest version for JEE Development)
- Status : Color is Verified

Last Updated: Dec 25 - 2022.

1. Prerequisites:

Option I

Start the VM using VM player and Logon to the server using telnet or directly in the VM console. Enter the root credentials to logon.

You can copy files from the host to VM using winscp.exe.

Option II.

Using docker:

Instantiate a container, kafkao.

You can copy files from the host to container using docker copy command.

Execute the following command, it will perform the following:

- Create a network : spark-net
- download the image, centos:7
- start a container with the name, kafkao and mount host folder in /opt

```
#docker network create --driver bridge spark-net
```

4 Kafka – Administration

```
#docker run --name kafkao --hostname kafkao -p 9092:9092 -p 8081:8081 -p 2181:2181 -p 9999:9999 -i -t --privileged --network spark-net -v /Users/henrypotsangbam/Documents/Software:/Software centos:8 /usr/sbin/init
```

Optional:

```
#docker run --name kafkao --hostname kafkao -p 9092:9092 -p 9081:8081 -p 9082:8082 -p 3181:2181 -p 9990:9999 -p 9021:9021 -i -t --privileged --network spark-net -v /Volumes/Samsung_T5/software:/Software -v /Volumes/Samsung_T5/software/install:/opt -v /Volumes/Samsung_T5/software/data:/data centos:7 /usr/sbin/init
```

You can verify the container from a separate terminal:

```
(base) Henrys-MacBook-Air:~ henrypotsangbam$ docker ps
CONTAINER ID        IMAGE               COMMAND             CREATED            STATUS              PORTS
0945477d457b        ehenry0227/learning:spark-kafka-raw   "/bin/bash"         34 seconds ago    Up 33 seconds
ds                  0.0.0.0:6062->6062/tcp, 0.0.0.0:8081->8081/tcp   kafka0
(base) Henrys-MacBook-Air:~ henrypotsangbam$
```

Perform the installation after this as a normal server.

Update the server.properties.

```
// Don't change the listeners address if the client also run in the host machine.
```

```
advertised.listeners=PLAINTEXT://localhost:9092
```

5 Kafka – Administration

After installation to start the zookeeper and kafka broker:

Open two separate terminals and execute the following

```
# /opt/kafka/bin/zookeeper-server-start.sh /opt/kafka/config/zookeeper.properties  
# /opt/kafka/bin/kafka-server-start.sh -daemon /opt/kafka/config/server.properties
```

Kafka 2nd Cluster: (It will be used as DC2 for Kafka Mirroring)

```
#docker run --name kafka1 --hostname kafka1 -p 7092:9092 -p 7081:8081 -p 4181:2181 -p  
8999:9999 -i -t --privileged --network spark-net -v  
/Volumes/Samsung_T5/software/:/Software -v  
/Volumes/Samsung_T5/software/install/:/opt -v  
/Volumes/Samsung_T5/software/data/:/data centos:7 /usr/sbin/init
```

Install the kafka as done before if required or copy the kafka folder in different name to persevere the libraries.

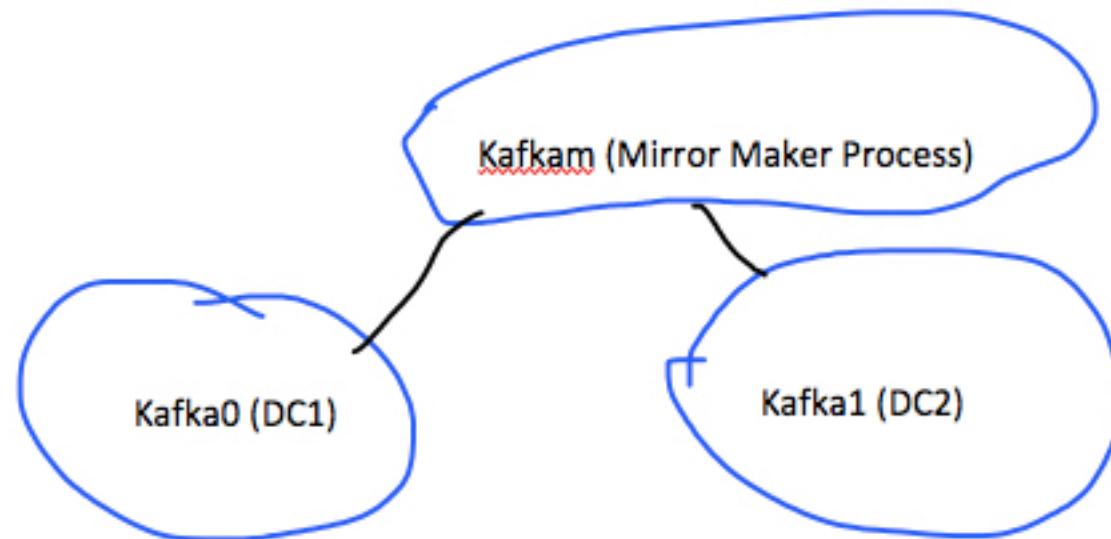
Kafka Mirror Maker: (It will execute the Kafka Mirroring process)

```
#docker run --name kafkam --hostname kafkam -p 6092:9092 -p 6081:8081 -i -t --  
privileged --network spark-net -v /Users/henrypotsangbam/Documents/Docker:/opt  
centos:7 /usr/sbin/init
```

Kafka Mirror Architecture:

6 Kafka – Administration

12:46 PM



Note:

If you are using docker ensure to update the server.properties with the following entries for accessing the broker from the host machine.

// Changes Begin

```
listeners=PLAINTEXT://0.0.0.0:9092,PLAINTEXT_HOST://0.0.0.0:8081  
advertised.listeners=PLAINTEXT://kafkao:9092,PLAINTEXT_HOST://localhost:8081  
listener.security.protocol.map=PLAINTEXT:PLAINTEXT,PLAINTEXT_HOST:PLAINTEXT
```

// Changes End

use container, kafkao to connect from Docker. However, use localhost:8081 for connecting from the Host machine.

[kafka_2.13-3.2.1.tgz](#)

jdk-11.0.16_linux-aarch64_bin.tar
apache-zookeeper-3.8.0-bin.tar.gz

2. Installation of Kafka – 90 Mins

You need to install java before installing zookeeper and Kafka.

Installation of Java

```
#tar -xvf jdk-11.0.16_linux* -C /opt  
#cd /opt  
#mv jdk* jdk
```

Set the above path in the PATH variable and JAVA_HOME

Update profile as follow:

```
#vi ~/.bashrc  
  
export JAVA_HOME=/opt/jdk  
export PATH=$PATH:$JAVA_HOME/bin
```

Update the shell scripts using the following command.

```
#bash
```

Installing Zookeeper

You can choose any of the options given below:

Option I (Fresh Installation): (We will use this for our lab)

The following steps install Zookeeper with a basic configuration in /opt/zookeeper.
Its configured to store data in /opt/data/zookeeper:

Extract the zookeeper archive file in /opt and rename the installation folder for brevity.

```
# tar -xvf apache-zookeeper* -C /opt  
#cd /opt  
#mv apache-zookeeper* /opt/zookeeper  
#mkdir -p /opt/data/zookeeper
```

Create a zookeeper configuration file and update with the following values.

```
#vi /opt/zookeeper/conf/zoo.cfg  
tickTime=2000  
dataDir=/opt/data/zookeeper  
clientPort=2181
```

Update the zoo.cfg with the above entries. You can save the file using esc+wq!

Start the zookeeper using the following scripts.

```
# /opt/zookeeper/bin/zkServer.sh start
```

```
[root@tos opt]# /opt/zookeeper/bin/zkServer.sh start
ZooKeeper JMX enabled by default
Using config: /opt/zookeeper/bin/.../conf/zoo.cfg
Starting zookeeper ... STARTED
[root@tos opt]#
```

Option II (Part of the Apache Kafka) – Skip.

```
#bin/zookeeper-server-start.sh config/zookeeper.properties
```

Option II Ends Here.

You can now validate that Zookeeper is running correctly in standalone mode by connecting to the client port and sending the four-letter command srvr:

```
#yum install telnet
```

```
#telnet localhost 2181
```

```
➤ srvr
```

11 Kafka – Administration

```
[root@tos opt]# telnet localhost 2181
Trying ::1...
Connected to localhost.
Escape character is '^].
srvr
Zookeeper version: 3.4.12-e5259e437540f349646870ea94dc2658c4e44b3b, built on 03/
27/2018 03:55 GMT
Latency min/avg/max: 0/0/0
Received: 2
Sent: 1
Connections: 1
Outstanding: 0
Zxid: 0x0
Mode: standalone
Node count: 4
Connection closed by foreign host.
[root@tos opt]#
```

After zookeeper installation, let us install the Kafka Broker.

Installation of Kafka Broker

The following example installs Kafka in /opt/kafka, configured to use the Zookeeper server started previously and to store the message log segments stored in /tmp/kafka-logs:

```
# tar -xvf kafka_2*-C /opt
#cd /opt
# mv kafka_2* /opt/kafka
# mkdir /opt/data/kafka-logs
```

Update the /opt/kafka/config/server.properties to store the Kafka Log in the above mention folder.

```
# The number of threads that the server uses for processing requests, which may include
# disk I/O
num.io.threads=8

# The send buffer (SO_SNDBUF) used by the socket server
socket.send.buffer.bytes=102400

# The receive buffer (SO_RCVBUF) used by the socket server
socket.receive.buffer.bytes=102400

# The maximum size of a request that the socket server will accept (protection against
# OOM)
socket.request.max.bytes=104857600

#####
##### Log Basics #####
#####

# A comma separated list of directories under which to store log files
log.dirs=/opt/data/kafka-logs

# The default number of log partitions per topic. More partitions allow greater
# parallelism for consumption, but this will also result in more files across
-- INSERT --
```

If you are using docker, kindly refer the prerequisite section for setting specific to docker.

Start the broker with the following command

```
# /opt/kafka/bin/kafka-server-start.sh -daemon /opt/kafka/config/server.properties
```

```
[root@tos opt]# /opt/kafka/bin/kafka-server-start.sh -daemon /opt/kafka/config/server.properties
[root@tos opt]# jps
3476 Kafka
3499 Jps
2895 QuorumPeerMain
[root@tos opt]#
```

```
#mkdir /opt/scripts
```

All the common execution scripts will be stored in the above folder.

The following scripts will start a zookeeper along with a broker. Create the following files and update with the following scripts. It will start the zookeeper and kafka broker using the mention script.

```
#cd /opt/scripts
#vi startABroker.sh
##### Scripts Begin #####
#!/usr/bin/env bash
```

```
# Start Zookeeper.  
/opt/zookeeper/bin/zkServer.sh start  
  
#Start Kafka Server  
/opt/kafka/bin/kafka-server-start.sh -daemon /opt/kafka/config/server.properties
```

```
echo "Started Successfully"
```

```
//##### Scripts End  
#####
```

To shutdown the Broker, find the process and kill.

```
ps -eaf | grep java
```

or

create a scripts : /opt/scripts/stopBroker.sh with the following commands in it.

```
#!/usr/bin/env bash
```

```
/opt/kafka/bin/kafka-server-stop.sh
```

To Stop Zookeeper create the following script. Don't include the ---- line.

```
#vi /opt/scripts/stopZookeeper.sh
```

Update the following commands in the above script and save it.

```
#!/usr/bin/env bash
```

```
# Stop Zookeeper.
```

```
/opt/zookeeper/bin/zkServer.sh stop  
echo "Stop zookeeper Successfully"
```

To reinitialize the Cluster.

```
#vi reinitializeCluster.sh
```

```
#!/usr/bin/env bash
```

```
rm -fr /opt/kafka/data/kafka-logs/*  
rm -fr /opt/kafka/data/zookeeper/*
```

```
cp /opt/kafka/config/server.properties_plain /opt/kafka/config/server.properties
```

```
echo "Reinitialize Successfully"
```

Lab Installation completes End here.

3. Installation Confluent Kafka (Local) – 30 Minutes

The purpose of this lab is to demonstrate the basic and most powerful capabilities of Confluent Platform – Schema Registry.

Install Confluent kafka with the following step.

Inflate the confluent kafka compress file as shown below:

```
#tar -xvf confluent-7.2.2.tar -C /opt
```

Rename the folder.

```
#cd /opt
```

```
# mv confluent-7.2.2 confluent
```

Set the environment variable for the Confluent Platform directory.

```
export CONFLUENT_HOME=/opt/confluent
```

Set your PATH variable:

```
# vi ~/.bashrc
```

```
export PATH=/opt/confluent/bin:${PATH};
```

type the following to activate the script.

```
#bash
```

----- Lab Ends Here -----

4. Basic Kafka Operations - CLI (Topic) – 30 Mins

In this lab you will be able to create a topic and perform some operations to understand the information about topic like partition and replication.

You need to start the broker using startABroker.sh. The script should be in /opt/scripts folder

```
#sh startABroker.sh  
#jps
```

```
[root@tos scripts]# sh startABroker.sh  
ZooKeeper JMX enabled by default  
Using config: /opt/zookeeper/bin/../conf/zoo.cfg  
Starting zookeeper ... STARTED  
Started Successfully  
[root@tos scripts]# jps  
11665 Jps  
11646 Kafka  
11375 QuorumPeerMain  
[root@tos scripts]#
```

Once the Kafka broker is started, we can verify that it is working by performing some simple operations against the broker; creating a test topic etc.

Create and verify details about topic:

```
# /opt/kafka/bin/kafka-topics.sh --create --bootstrap-server kafka0:9092 --replication-factor 1 --partitions 1 --topic test
```

```
# /opt/kafka/bin/kafka-topics.sh --create --bootstrap-server kafka0:9092 --replication-factor 1 --partitions 12 --topic IBM  
  
# /opt/kafka/bin/kafka-topics.sh --bootstrap-server kafka0:9092 --describe --topic test  
  
#/opt/kafka/bin/kafka-topics.sh --bootstrap-server kafka0:9092 --describe --topic IBM
```

Verify the no of partition in the output.

list and describe topic.

```
# /opt/kafka/bin/kafka-topics.sh --bootstrap-server kafka0:9092 --list
```

List and describe Topics

What does the tool do?

This tool lists the information for a given list of topics. If no topics are provided in the command line, the tool queries zookeeper to get all the topics and lists the information for them. The fields that the tool displays are - topic name, partition, leader, replicas, isr.

How to use the tool?

List only single topic named "test" (prints only topic name)

```
#/opt/kafka/bin/kafka-topics.sh --list --bootstrap-server kafka:9092 --topic test
```

List all topics (prints only topic names)

```
#/opt/kafka/bin/kafka-topics.sh --list --bootstrap-server kafka:9092
```

Describe only single topic named "test" (prints details about the topic)

```
#/opt/kafka/bin/kafka-topics.sh --describe --bootstrap-server kafka:9092 --topic test
```

Describe all topics (prints details about the topics)

```
#/opt/kafka/bin/kafka-topics.sh --describe --bootstrap-server kafka:9092
```

We will understand the output in details later.

Create Topics

What does the tool do?

By default, Kafka auto creates topic if "auto.create.topics.enable" is set to true on the server. This creates a topic with a default number of partitions, replication factor and uses Kafka's default scheme to do replica assignment. Sometimes, it may be required that we would like

to customize a topic while creating it. This tool helps to create a topic and also specify the number of partitions, replication factor and replica assignment list for the topic.

How to use the tool?

create topic with default settings

```
#/opt/kafka/bin/kafka-topics.sh --create --bootstrap-server kafka0:9092 --topic topic1 --partitions 2 --replication-factor 1
```

Create a topic with replication 2.

```
#/opt/kafka/bin/kafka-topics.sh --create --bootstrap-server kafka0:9092 --topic topic2 --partitions 2 --replication-factor 2
```

As shown above, it generates an error. Since there is only a single broker. It will fix later.

Lab CLI completes End here.

5. Basic Kafka Operations - CLI (Producer & Consumer) – 30 Mins

In this lab we will send message to the broker and consumer message using the kafka inbuilt commands.

You need to complete the previous lab before proceeding ahead.

You need to start the broker using startABroker.sh if not done earlier. The script should be in /opt/scripts folder

```
#sh startABroker.sh  
#jps
```

```
[root@tos scripts]# sh startABroker.sh  
ZooKeeper JMX enabled by default  
Using config: /opt/zookeeper/bin/..../conf/zoo.cfg  
Starting zookeeper ... STARTED  
Started Successfully  
[root@tos scripts]# jps  
11665 Jps  
11646 Kafka  
11375 QuorumPeerMain  
[root@tos scripts]#
```

Sent message to **test** topic: Open a console to send message to the topic, test. Enter some text as shown below.

```
# /opt/kafka/bin/kafka-console-producer.sh --broker-list kafka:9092 --topic test
Test Message 1
Test Message 2
^D
#
```

Consume messages from a test topic: As soon as you enter the following script in a separate terminal, you should be able to consume the messages that we have type in the producer console.

```
# /opt/kafka/bin/kafka-console-consumer.sh --bootstrap-server kafka:9092 --topic test --
from-beginning
```

Lab CLI ends here.

6. Zookeeper – 120 Minutes

In this lab, we will perform familiarization of ZK cli features and configure zookeeper ensemble:

To perform ZooKeeper CLI operations, first start your ZooKeeper server and then, ZooKeeper client by executing “bin/zkCli.sh”, which is in the Broker installation folder.

/opt/zookeeper/bin

#bin/zkCli.sh

```
zookeeper -- java - sudo — 137x24
2020-08-21 12:40:05,342 [myid:] - INFO  [main:Environment@98] - Client environment:os.memory.free=116MB
2020-08-21 12:40:05,345 [myid:] - INFO  [main:Environment@98] - Client environment:os.memory.max=228MB
2020-08-21 12:40:05,345 [myid:] - INFO  [main:Environment@98] - Client environment:os.memory.total=123MB
2020-08-21 12:40:05,355 [myid:] - INFO  [main:ZooKeeper@1005] - Initiating client connection, connectString=localhost:2181 sessionTimeout=30000 watcher=org.apache.zookeeper.ZooKeeperMain$MyWatcher@4926097b
2020-08-21 12:40:05,361 [myid:] - INFO  [main:X509Util@77] - Setting -D jdk.tls.rejectClientInitiatedRenegotiation=true to disable client-initiated TLS renegotiation
2020-08-21 12:40:05,378 [myid:] - INFO  [main:ClientCnxnSocket@239] - jute.maxbuffer value is 1048575 Bytes
2020-08-21 12:40:05,389 [myid:] - INFO  [main:ClientCnxn@1703] - zookeeper.request.timeout value is 0. feature.enabled=false
Welcome to ZooKeeper!
2020-08-21 12:40:05,401 [myid:localhost:2181] - INFO  [main-SendThread(localhost:2181):ClientCnxn$SendThread@1154] - Opening socket connection to server localhost/127.0.0.1:2181.
2020-08-21 12:40:05,407 [myid:localhost:2181] - INFO  [main-SendThread(localhost:2181):ClientCnxn$SendThread@1156] - SASL config status: Will not attempt to authenticate using SASL (unknown error)
JLine support is enabled
2020-08-21 12:40:05,505 [myid:localhost:2181] - INFO  [main-SendThread(localhost:2181):ClientCnxn$SendThread@986] - Socket connection established, initiating session, client: /127.0.0.1:50738, server: localhost/127.0.0.1:2181
2020-08-21 12:40:05,517 [myid:localhost:2181] - INFO  [main-SendThread(localhost:2181):ClientCnxn$SendThread@1420] - Session establishment complete on server localhost/127.0.0.1:2181, session id = 0x1000075b4060002, negotiated timeout = 30000
WATCHER::

WatchedEvent state:SyncConnected type:None path:null
[zk: localhost:2181(CONNECTED) 0]
```

or New CLI

```
[root@kafka0 kafka]# bin/zookeeper-shell.sh localhost:2181
Connecting to localhost:2181
Welcome to ZooKeeper!
JLine support is disabled

WATCHER::

WatchedEvent state:SyncConnected type:None path:null

ls /
[admin, brokers, cluster, config, consumers, controller, controller_epoch, feature, isr_change_notification, l
atest_producer_id_block, log_dir_event_notification, zookeeper]
ls -R /
/
/admin
/brokers
/cluster
```

As highlighted above, the cli is connected to the zookeeper running on localhost listening at port 2181.

Let us get all znode:

```
#ls /
```

```
[zk: localhost:2181(CONNECTED) 1] ls /
[admin, brokers, cluster, config, consumers, controller, controller_epoch, isr_change_notification, latest_producer_id_block, log_dir_eve
nt_notification, zookeeper]
[zk: localhost:2181(CONNECTED) 2]

#ls -R /
#ls -R /brokers
#get /brokers/ids/o

[[zk: localhost:2181(CONNECTED) 22] get /brokers/ids/0
{"listener_security_protocol_map":{"PLAINTEXT":"PLAINTEXT"},"endpoints":["PLAINTEXT://192.168.0.111:9092"],"jmx_port":-1,"host":"192.168.
0.111","timestamp":1597992632471,"port":9092,"version":4}
[zk: localhost:2181(CONNECTED) 23]
```

Information about the broker which is register in the znode.

Here, the Node 0 is own by the broker IP – 192.168.0.111 with the port 9092.

This way you can determine the Node details whenever there are any issues as shown below:

```
[2020-08-21 12:52:15,093] WARN [Consumer clientId=consumer-console-consumer-26076-1, groupId=console-consumer
-26076] Connection to node -1 (localhost/127.0.0.1:9020) could not be established. Broker may not be availabl
e. (org.apache.kafka.clients.NetworkClient)
[2020-08-21 12:52:15,093] WARN [Consumer clientId=consumer-console-consumer-26076-1, groupId=console-consumer
-26076] Bootstrap broker localhost:9020 (id: -1 rack: null) disconnected (org.apache.kafka.clients.NetworkCli
ent)
^CProcessed a total of 0 messages
(base) Henrys-MacBook-Air:zookeeper henrynottingham$
```

Using the Kafka tool `zookeeper-shell.sh` (Broker Installation folder) we can connect to a ZooKeeper host in our cluster and look at how data is stored.

```
#zookeeper-shell.sh localhost:2181  
#ls /brokers/topics
```

If we look at the path /brokers/topics you should see a list of the topics that you have created.

```
[(base) Henrys-MacBook-Air:bin henrypotsangbam$ sh zookeeper-shell.sh localhost:2181  
Connecting to localhost:2181  
Welcome to ZooKeeper!  
JLine support is disabled  
  
WATCHER::  
  
WatchedEvent state:SyncConnected type:None path:null  
  
ls  
ls [-s] [-w] [-R] path  
ls /brokers/topics  
[__consumer_offsets, quickstart-events, test, topic1]
```

You should also be able to see the topic __consumer_offsets. That topic is one that you did not create, it is in fact a private topic used internal by Kafka itself. This topic stored the committed offsets for each topic and partition per group id.

The path /controller exists in zookeeper and we are running one command to look at that current value

```
zookeeper-shell.sh localhost:2181  
get /controller
```

Now in the next section, let us configure cluster of zookeeper – 3 nodes.
Ensemble

Setting up a ZooKeeper Ensemble

Create 3 zookeepers' configuration to create zookeeper ensemble of 3 nodes:

The <ZOOKEEPER_HOME>/conf/zoo1.cfg file should have the content:

```
dataDir=/opt/data/zookeeper/1  
clientPort=2181  
initLimit=5  
syncLimit=2  
server.1=localhost:2888:3888  
server.2=localhost:2889:3889  
server.3=localhost:2890:3890
```

The <ZOOKEEPER_HOME>/conf/zoo2.cfg file should have the content:

```
dataDir=/opt/data/zookeeper/2
clientPort=2182
initLimit=5
syncLimit=2
server.1=localhost:2888:3888
server.2=localhost:2889:3889
server.3=localhost:2890:3890
```

The <ZOOKEEPER_HOME>/conf/zoo3.cfg file should have the content:

```
dataDir=/opt/data/zookeeper/3
clientPort=2183
initLimit=5
syncLimit=2
server.1=localhost:2888:3888
server.2=localhost:2889:3889
server.3=localhost:2890:3890
```

To complete your multi-node configuration, you will specify a node ID on each of the servers. To do this, you will create a myid file on each node. Each file will contain a number that correlates to the server number assigned in the configuration file.

```
#cd /opt/data/zookeeper/  
#mkdir 1 2 3  
#vi 1/myid → 1  
2/myid → 2  
3/myid → 3
```

```
[(base) Henrys-MacBook-Air:zookeeper henrypotsangbam$ sudo vi 1/myid  
[(base) Henrys-MacBook-Air:zookeeper henrypotsangbam$ sudo vi 2/myid  
[(base) Henrys-MacBook-Air:zookeeper henrypotsangbam$ sudo vi 3/myid  
[(base) Henrys-MacBook-Air:zookeeper henrypotsangbam$ more 1/myid  
1  
[(base) Henrys-MacBook-Air:zookeeper henrypotsangbam$ pwd  
/opt/data/zookeeper  
(base) Henrys-MacBook-Air:zookeeper henrypotsangbam$ ]
```

To start the servers, you can simply explicitly reference the configuration files:

```
cd <ZOOKEEPER_HOME>  
bin/zkServer.sh start conf/zoo1.cfg  
bin/zkServer.sh start conf/zoo2.cfg  
bin/zkServer.sh start conf/zoo3.cfg
```

```
*@* QuorumPeerMain
[(base) Henrys-MacBook-Air:zookeeper henrypotsangbam$ sudo bin/zkServer.sh start conf/zoo2.cfg
/usr/bin/java
ZooKeeper JMX enabled by default
Using config: conf/zoo2.cfg
Starting zookeeper ... STARTED
[(base) Henrys-MacBook-Air:zookeeper henrypotsangbam$ sudo bin/zkServer.sh start conf/zoo3.cfg
/usr/bin/java
ZooKeeper JMX enabled by default
Using config: conf/zoo3.cfg
Starting zookeeper ... STARTED
[(base) Henrys-MacBook-Air:zookeeper henrypotsangbam$ jps
4136 Jps
[(base) Henrys-MacBook-Air:zookeeper henrypotsangbam$ sudo jps
4131 QuorumPeerMain
4139 Jps
4107 QuorumPeerMain
4078 QuorumPeerMain
(base) Henrys-MacBook-Air:zookeeper henrypotsangbam$
```

You will now start a ZooKeeper command line client and connect to ZooKeeper on **node 1**:

```
#bin/zkCli.sh -server localhost:2181
```

List the newly created znode:

```
#ls /
```

You can verify the leader using command:

```
#cd /opt/zookeeper  
#bin/zkServer.sh status conf/zoo1.cfg  
# bin/zkServer.sh status conf/zoo2.cfg  
# bin/zkServer.sh status conf/zoo3.cfg
```

```
[root@kafka0 zookeeper]# bin/zkServer.sh status conf/zoo2.cfg  
ZooKeeper JMX enabled by default  
Using config: conf/zoo2.cfg  
Client port found: 2182. Client address: localhost.  
Mode: leader  
[root@kafka0 zookeeper]# bin/zkServer.sh status conf/zoo1.cfg  
ZooKeeper JMX enabled by default  
Using config: conf/zoo1.cfg  
Client port found: 2181. Client address: localhost.  
Mode: follower  
[root@kafka0 zookeeper]# bin/zkServer.sh status conf/zoo3.cfg  
ZooKeeper JMX enabled by default  
Using config: conf/zoo3.cfg  
Client port found: 2183. Client address: localhost.  
Mode: follower  
[root@kafka0 zookeeper]# pwd  
/opt/zookeeper
```

Here, the leader is running in 2182 port.

Finally, let us configure the broker to point to the zookeeper ensemble.

The configuration will be in the Kafka conf folder.

Copy conf/server.properties into conf/server1.properties and modify the following two properties in it. Use cp command.

```
config — vi server.properties — 80x24

# The maximum size of a log segment file. When this size is reached a new log se
gment will be created.
log.segment.bytes=1073741824

# The interval at which log segments are checked to see if they can be deleted a
ccording
# to the retention policies
log.retention.check.interval.ms=300000

#####
# Zookeeper connection string (see zookeeper docs for details).
# This is a comma separated host:port pairs, each corresponding to a zk
# server. e.g. "127.0.0.1:3000,127.0.0.1:3001,127.0.0.1:3002".
# You can also append an optional chroot string to the urls to specify the
# directory, e.g. "/kafka".
zookeeper.connect=localhost:2181,localhost:2182,localhost:2183

# Timeout in ms for connecting to zookeeper
zookeeper.connection.timeout.ms=18000

@
"server.properties" 136L, 6884C written
```

log.dirs=/opt/data/kafka-logs1

Updated the Zookeeper.connect property to connect to all zookeeper nodes.

Start all the zookeepers if not done.

```
#cd /opt/zookeeper
```

```
[(base) Henrys-MacBook-Air:zookeeper henrypotsangbam$ sudo bin/zkServer.sh start ]  
conf/zoo1.cfg  
[Password:  
/usr/bin/java  
ZooKeeper JMX enabled by default  
Using config: conf/zoo1.cfg  
Starting zookeeper ... STARTED  
[(base) Henrys-MacBook-Air:zookeeper henrypotsangbam$ sudo bin/zkServer.sh start ]  
conf/zoo2.cfg  
/usr/bin/java  
ZooKeeper JMX enabled by default  
Using config: conf/zoo2.cfg  
Starting zookeeper ... STARTED  
[(base) Henrys-MacBook-Air:zookeeper henrypotsangbam$ sudo bin/zkServer.sh start ]  
conf/zoo3.cfg  
/usr/bin/java  
ZooKeeper JMX enabled by default  
Using config: conf/zoo3.cfg  
Starting zookeeper ... STARTED  
(base) Henrys-MacBook-Air:zookeeper henrypotsangbam$ █
```

Start the broker using the modified configuration.

```
#/opt/kafka/bin/kafka-server-start.sh -daemon /opt/kafka/config/server1.properties
```

You can verify from the server.log that the broker is referring to the new property as shown below: /opt/kafka/logs/server.log

```
[2020-08-24 08:41:05,479] INFO starting (kafka.server.KafkaServer)
[2020-08-24 08:41:05,482] INFO Connecting to zookeeper on localhost:2181,localhost:2182,localhost:2183 (kafka.server.KafkaServer)
[2020-08-24 08:41:05,538] INFO [ZooKeeperClient Kafka server] Initializing a new session to localhost:2181,localhost:2182,localhost:2183. (kafka.zookeeper.ZooKeeperClient)
[2020-08-24 08:41:05,564] INFO Client environment:zookeeper.version=3.5.7-f0fdd52973d373ffd9c86b81d99842dc2c7f660e
, built on 02/10/2020 11:30 GMT (org.apache.zookeeper.ZooKeeper)

[2020-08-24 08:47:42,893] INFO Kafka version: 2.5.0 (org.apache.kafka.common.utils.AppInfoParser)
[2020-08-24 08:47:42,893] INFO Kafka commitId: 66563e712b0b9f84 (org.apache.kafka.common.utils.AppInfoParser)
[2020-08-24 08:47:42,895] INFO Kafka startTimeMs: 1599229042881 (org.apache.kafka.common.utils.AppInfoParser)
[2020-08-24 08:47:42,902] INFO [KafkaServer id=0] started (kafka.server.KafkaServer)
(base) Henrys-MacBook-Air:~ henry$
```

Let us perform some activities using the new set up.

```
# /opt/kafka/bin/kafka-topics.sh --create --bootstrap-server localhost:9092 --replication-factor 1 --
partitions 2 --topic testz
```

```
# /opt/kafka/bin/kafka-topics.sh --bootstrap-server localhost:9092 --describe --topic testz
```

```
# /opt/kafka/bin/kafka-console-producer.sh --broker-list localhost:9092 --topic testz
```

Test Message 1

Test Message 2

To exit : ^D

Consume the above messages:

```
# /opt/kafka/bin/kafka-console-consumer.sh --bootstrap-server localhost:9092 --topic testz --from-beginning
```

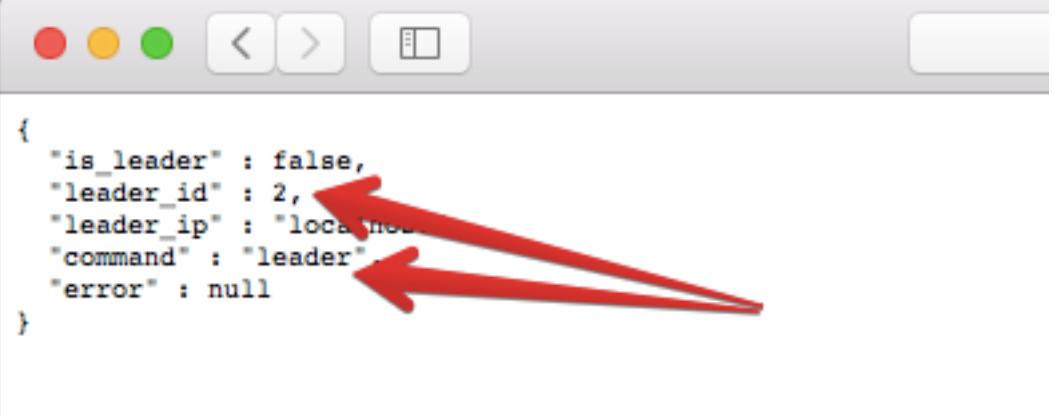
```
((base) Henrys-MacBook-Air:logs henrypotsangbam$ sudo /opt/kafka/bin/kafka-console-producer.sh --broker-list localhost:9092 --topic testz
>Test Message 1
Test Message 2
[>>(base) Henrys-MacBook-Air:logs henrypotsangbam$ sudo /opt/kafka/bin/kafka-console-consumer.sh --bootstrap-server localhost:9092 --topic testz --from-beginning
Test Message 1
Test Message 2
```

Let us verify the Leader of the zookeeper and Kill the leader process.

<http://localhost:8080/commands/leader>

[or CLI :](#)

```
#/opt/zookeeper/bin/zkServer.sh status /opt/zookeeper/conf/zoo2.cfg
```



```
{  
    "is_leader" : false,  
    "leader_id" : 2,  
    "leader_ip" : "localhost",  
    "command" : "leader",  
    "error" : null  
}
```

```
[root@kafka0 /]# /opt/zookeeper/bin/zkServer.sh status /opt/zookeeper/conf/zoo2.cfg  
ZooKeeper JMX enabled by default  
Using config: /opt/zookeeper/conf/zoo2.cfg  
Client port found: 2182. Client address: localhost.  
Mode: leader  
[root@kafka0 /]#
```

You can verify the other configuration if 2 is not the leader.

Since 2 is the leader. Let us kill the process Id of the Zookeeper id 2.

```
#ps -eaf | grep zoo2.cfg
```

```
[base] Henrys-MacBook-Air:config henrypotsangbam$ sudo ps -eaf | grep zoo2.cfg
 501  3725  537  0  9:05AM ttys000  0:00.00 grep zoo2.cfg
  0  684   1  0  8:39AM ttys001  0:05.91 /usr/bin/java -Dzookeeper.log.dir=/opt/zookeeper/bin/../logs -Dzookeeper.log.fil
e=zookeeper-root-server-Henrys-MacBook-Air.local.log -Dzookeeper.root.logger=INFO,CONSOLE -XX:+HeapDumpOnOutOfMemoryError -XX:OnO
utOfMemoryError=kill -9 %p -cp /opt/zookeeper/bin/../zookeeper-server/target/classes:/opt/zookeeper/bin/../build/classes:/opt/zoo
keeper/bin/../zookeeper-server/target/lib/*.jar:/opt/zookeeper/bin/../build/lib/*.jar:/opt/zookeeper/bin/../lib/zookeeper-prometh
eus-metrics-3.6.1.jar:/opt/zookeeper/bin/../lib/zookeeper-jute-3.6.1.jar:/opt/zookeeper/bin/../lib/zookeeper-3.6.1.jar:/opt/zooke
eper/bin/../lib/snappy-java-1.1.7.jar:/opt/zookeeper/bin/../lib/slf4j-log4j12-1.7.25.jar:/opt/zookeeper/bin/../lib/slf4j-api-1.7.
25.jar:/opt/zookeeper/bin/../lib/simpleclient_servlet-0.6.0.jar:/opt/zookeeper/bin/../lib/simpleclient_hotspot-0.6.0.jar:/opt/zoo
keeper/bin/../lib/simpleclient_common-0.6.0.jar:/opt/zookeeper/bin/../lib/simpleclient-0.6.0.jar:/opt/zookeeper/bin/../lib/netty-
transport-native-unix-common-4.1.48.Final.jar:/opt/zookeeper/bin/../lib/netty-transport-native-epoll-4.1.48.Final.jar:/opt/zooke
per/bin/../lib/netty-transport-4.1.48.Final.jar:/opt/zookeeper/bin/../lib/netty-resolver-4.1.48.Final.jar:/opt/zookeeper/bin/..
1ib/netty-handler-4.1.48.Final.jar:/opt/zookeeper/bin/../lib/netty-common-4.1.48.Final.jar:/opt/zookeeper/bin/../lib/netty-codec-4
.1.48.Final.jar:/opt/zookeeper/bin/../lib/netty-buffer-4.1.48.Final.jar:/opt/zookeeper/bin/../lib/metrics-core-3.2.5.jar:/opt/zoo
keeper/bin/../lib/log4j-1.2.17.jar:/opt/zookeeper/bin/../lib/json-simple-1.1.1.jar:/opt/zookeeper/bin/../lib/jline-2.11.jar:/opt/
zookeeper/bin/../lib/jetty-util-9.4.24.v20191120.jar:/opt/zookeeper/bin/../lib/jetty-servlet-9.4.24.v20191120.jar:/opt/zookeeper/
bin/../lib/jetty-server-9.4.24.v20191120.jar:/opt/zookeeper/bin/../lib/jetty-security-9.4.24.v20191120.jar:/opt/zookeeper/bin/..
lib/jetty-io-9.4.24.v20191120.jar:/opt/zookeeper/bin/../lib/jetty-http-9.4.24.v20191120.jar:/opt/zookeeper/bin/../lib/javax.servl
et-api-3.1.0.jar:/opt/zookeeper/bin/../lib/jackson-databind-2.10.3.jar:/opt/zookeeper/bin/../lib/jackson-core-2.10.3.jar:/opt/zoo
keeper/bin/../lib/jackson-annotations-2.10.3.jar:/opt/zookeeper/bin/../lib/commons-lang-2.6.jar:/opt/zookeeper/bin/../lib/commons
-clli-1.2.jar:/opt/zookeeper/bin/../lib/audience-annotations-0.5.0.jar:/opt/zookeeper/bin/..zoo*er-*.jar:/opt/zookeeper/bin/..
/zookeeper-server/src/main/resources/lib/*.jar:/opt/zookeeper/bin/..conf: -Xmx1000m -com.sun.management.jmxremote -Dcom.sun.man
agement.jmxremote.local.only=false org.apache.zookeeper.server.quorum.QuorumPeerMain conf/zoo2.cfg
(base) Henrys-MacBook-Air:config henrypotsangbam$
```

In this case, the process id is 684.

```
# kill -9 684
```

```
-----, -----
(base) Henrys-MacBook-Air:config henrypotsangbam$ sudo kill -9 684
(base) Henrys-MacBook-Air:config henrypotsangbam$ sudo jps
3745 Jps
659 QuorumPeerMain
3382 ConsoleConsumer
710 QuorumPeerMain
2026 Kafka
(base) Henrys-MacBook-Air:config henrypotsangbam$
```

You can refresh the browser to determine the new leader.

```
{  
    "is_leader" : false,  
    "leader_id" : 3,  
    "leader_ip" : "localhost",  
    "command" : "leader",  
    "error" : null  
}
```

or using cli.

```
[root@kafka0 /]# /opt/zookeeper/bin/zkServer.sh status /opt/zookeeper/conf/zoo2.cfg  
ZooKeeper JMX enabled by default  
Using config: /opt/zookeeper/conf/zoo2.cfg  
Client port found: 2182. Client address: localhost.  
Error contacting service. It is probably not running.  
[root@kafka0 /]# /opt/zookeeper/bin/zkServer.sh status /opt/zookeeper/conf/zoo3.cfg  
ZooKeeper JMX enabled by default  
Using config: /opt/zookeeper/conf/zoo3.cfg  
Client port found: 2183. Client address: localhost.  
Mode: leader
```

So which node is the leader in your case? Here its 3

Try sending some messages again:

```
(base) Henrys-MacBook-Air:logs henrypotsangbam$ sudo /opt/kafka/bin/kafka-console-producer.sh --broker-list localhost:9092 --topic testz
Password:
>Another Message
>
^CProcessed a total of 0 messages
(base) Henrys-MacBook-Air:config henrypotsangbam$ sudo /opt/kafka/bin/kafka-console-consumer.sh --bootstrap-server localhost: 9092 --topic testz --from-beginning
Test Message 1
Test Message 2
Another Message
```

You should be able to consume messages without any issue.

Get some config details from the zookeeper:

```
get /controller
get /brokers/topics/testz
get /brokers/ids/0
get /zookeeper/config
```

```
[[zk: localhost:2181(CONNECTED) 1] get /controller
{"version":1,"brokerid":0,"timestamp":"1598239062599"}
[[zk: localhost:2181(CONNECTED) 20] get /brokers/topics/testz
{"version":2,"partitions": {"1": [0], "0": [0]}, "adding_replicas": {}, "removing_replicas": {}}
[[zk: localhost:2181(CONNECTED) 21] get /brokers/ids/0
 {"listener_security_protocol_map": {"PLAINTEXT": "PLAINTEXT"}, "endpoints": ["PLAINTEXT://quickstart.cloudera:9092"], "jmx_port": -1, "host": "quickstart.cloudera", "timestamp": "1598239062331", "port": 9092, "version": 4}
[[zk: localhost:2181(CONNECTED) 22] get /zookeeper/config
server.1=localhost:2888:participant
server.2=localhost:2889:participant
server.3=localhost:2890:participant
version=0
[zk: localhost:2181(CONNECTED) 23]
```

----- lab Ends Here -----

Log:

server.1=localhost:2888:3888 (2181)
server.2=localhost:2889:3889. (2182) - First Leader
server.3=localhost:2890:3890 -(2183) Second Leader

Scenarios : 2 was leader , after 2 shutdown -> 3 Became a leader

7. Kafka cluster – 90 Minutes

Understanding Kafka Failover in a cluster environment.

In this tutorial, we are going to run many Kafka Nodes on our development laptop so, you will need at least 8 GB of RAM for local dev machine. You can run just two servers if you have less memory than 8 GB.

We are going to create a replicated topic and then demonstrate consumer along with broker failover. We also demonstrate load balancing of Kafka consumers.

We show how, with many groups, Kafka acts like a Publish/Subscribe. But, when we put all of our consumers in the same group, Kafka will load share the messages to the consumers in the same group (more like a queue than a topic in a traditional MOM sense).

If not already running, start ZooKeeper.

Also, shut down Kafka from the first tutorial.

Next, you need to copy server properties for three brokers (detailed instructions to follow). Then we will modify these Kafka server properties to add unique Kafka ports, Kafka log locations, and unique Broker ids. Then we will create three scripts to start these servers up using these properties, and then start the servers.

Lastly, we create replicated topic and use it to demonstrate Kafka consumer failover, and Kafka broker failover.

Create three new Kafka server-n.properties files

In this section, we will copy the existing Kafka `server.properties` to `server-0.properties`, `server-1.properties`, and `server-2.properties`.

Then change `server-0.properties` to set `log.dirs` to “`/opt/data/kafka-logs/kafka-0`. Then we modify `server-1.properties` to set `port` to `9093`, broker `id` to `1`, and `log.dirs` to “`/opt/data/kafka-logs/kafka-1`”.

Lastly modify `server-2.properties` to use `port` `9094`, broker `id` `2`, and `log.dirs` “`/opt/data/kafka-logs/kafka-2`”.

Copy server properties file as follows: We will store all server's configuration in a single folder config.

```
$ cd /opt  
$ mkdir -p kafka-config/config  
$ cp kafka/config/server.properties kafka-config/config/server-0.properties  
$ cp kafka/config/server.properties kafka-config/config/server-1.properties  
$ cp kafka/config/server.properties kafka-config/config/server-2.properties
```

```
[root@tos opt]# mkdir -p kafka-config/config
[root@tos opt]# pwd
/opt
[root@tos opt]# ls
couchbase  data  jdk1.8.0_45  kafka  kafka-config  zookeeper  zookeeper.out
[root@tos opt]# cp kafka/config/server.properties kafka-config/server-0.properties
[root@tos opt]# ls
couchbase  data  jdk1.8.0_45  kafka  kafka-config  zookeeper  zookeeper.out
[root@tos opt]# cp kafka/config/server.properties kafka-config/config/server-1.properties
[root@tos opt]# cp kafka/config/server.properties kafka-config/config/server-2.properties
[root@tos opt]# █
```

With your favourite text editor update server-0.properties so that `log.dirs` is set to `./logs/kafka-0`. Leave the rest of the file the same. Make sure `log.dirs` is only defined once.

```
#vi /opt/kafka-config/config/server-0.properties
broker.id=0
listeners=PLAINTEXT://kafkao:9092
advertised.listeners=PLAINTEXT://kafkao:9092
log.dirs=/opt/data/kafka-logs/kafka-0
```

With your favorite text editor change `log.dirs`, `broker.id` and `log.dirs` of `server-1.properties` as follows.

```
#vi /opt/kafka-config/config/server-1.properties
```

```
broker.id=1
listeners=PLAINTEXT://kafkao:9093
advertised.listeners=PLAINTEXT://kafkao:9093
log.dirs=/opt/data/kafka-logs/kafka-1
```

With your favorite text editor change `log.dirs`, `broker.id` and `log.dirs` of `server-2.properties` as follows.

```
#vi /opt/kafka-config/config/server-2.properties
```

```
broker.id=2
listeners=PLAINTEXT://kafkao:9094
advertised.listeners=PLAINTEXT://kafkao:9094
log.dirs=/opt/data/kafka-logs/kafka-2
```

Create Startup scripts for these three Kafka servers

The startup scripts will just run `kafka-server-start.sh` with the corresponding properties file.

```
#vi /opt/kafka-config/start-1st-server.sh
#!/usr/bin/env bash
## Run Kafka
/opt/kafka/bin/kafka-server-start.sh "/opt/kafka-config/config/server-0.properties"
```

```
#vi /opt/kafka-config/start-2nd-server.sh
#!/usr/bin/env bash
## Run Kafka
/opt/kafka/bin/kafka-server-start.sh "/opt/kafka-config/config/server-1.properties"
```

```
#vi /opt/kafka-config/start-3rd-server.sh
#!/usr/bin/env bash
## Run Kafka
/opt/kafka/bin/kafka-server-start.sh "/opt/kafka-config/config/server-2.properties"
```

Notice we are passing the Kafka server properties files that we created in the last step.
Now run all three in separate terminals/shells.

Run Kafka servers each in own terminal from /opt/kafka-config

```
#cd /opt/kafka-config
// to make the scripts executable.
#chmod 755 start-1st-server.sh start-2nd-server.sh start-3rd-server.sh
$ ./start-1st-server.sh

$ ./start-2nd-server.sh

$ ./start-3rd-server.sh
```

Give these servers a couple of minutes to startup and connect to ZooKeeper.

```
[root@tos:/opt/kafka-config] [2018-05-16 15:17:50,404] INFO [GroupCoordinator 2]: Starting up. (kafka.coordinator.group.GroupCoordinator)
[2018-05-16 15:17:50,411] INFO [GroupCoordinator 2]: Startup complete. (kafka.coordinator.group.GroupCoordinator)
[2018-05-16 15:17:50,564] INFO [GroupMetadataManager brokerId=2] Removed 0 expired offsets in 140 milliseconds. (kafka.coordinator.group.GroupMetadataManager)
[2018-05-16 15:17:50,637] INFO [ProducerId Manager 2]: Acquired new producerId block (brokerId:2,blockStartProducerId:4000,blockEndProducerId:4999) by writing to Zk with path version 5 (kafka.coordinator.transaction.ProducerIdManager)
[2018-05-16 15:17:50,777] INFO [TransactionCoordinator id=2] Starting up. (kafka.coordinator.transaction.TransactionCoordinator)
[2018-05-16 15:17:50,788] INFO [TransactionCoordinator id=2] Startup complete. (kafka.coordinator.transaction.TransactionCoordinator)
[2018-05-16 15:17:51,303] INFO [Transaction Marker Channel Manager 2]: Starting (kafka.coordinator.transaction.TransactionMarkerChannelManager)
[2018-05-16 15:17:51,574] INFO [/config/changes-event-process-thread]: Starting (kafka.common.ZkNodeChangeNotificationListener$ChangeEventProcessThread)
[2018-05-16 15:17:51,723] INFO Kafka version : 1.1.0 (org.apache.kafka.common.utils.AppInfoParser)
[2018-05-16 15:17:51,723] INFO Kafka commitId : fdcf75ea326b8e07 (org.apache.kafka.common.utils.AppInfoParser)
[2018-05-16 15:17:51,728] INFO [KafkaServer id=2] started (kafka.server.KafkaServer)
[2018-05-16 15:15:08,676] INFO [Transaction Marker Channel Manager 1]: Starting (kafka.coordinator.transaction.TransactionMarkerChannelManager)
[2018-05-16 15:15:09,267] INFO [/config/changes-event-process-thread]: Starting (kafka.common.ZkNodeChangeNotificationListener$ChangeEventProcessThread)
Give thes[2018-05-16 15:15:09,281] INFO Kafka version : 1.1.0 (org.apache.kafka.common.utils.AppInfoParser)
[2018-05-16 15:15:09,282] INFO Kafka commitId : fdcf75ea326b8e07 (org.apache.kafka.common.utils.AppInfoParser)
[2018-05-16 15:15:09,285] INFO [KafkaServer id=1] started (kafka.server.KafkaServer)

@tos:/opt/kafka-config [2018-05-16 15:09:28,919] INFO [Log partition=test-0, dir=/opt/data/kafka-logs/kafka-0] Completed load of log with 1 segments, log start offset 0 and log end offset 154 ms (kafka.log.Log)
[2018-05-16 15:09:28,940] INFO Created log for partition test-0 in /opt/data/kafka-0 with properties {compression.type->producer, message.format.version->1.1-IV0, file.delete.delay.ms->60000, max.message.bytes->1000012, message.lag.ms->0, message.timestamp.type->CreateTime, min.insync.replicas->1, segment.jitter.ms->0, preallocate->false, min.cleanable.dirty.ratio->0.5, index.interval.bytes->4096, unclean.leader.election.enable->false}
[2018-05-16 15:09:28,940] INFO [Log partition=test-0, dir=/opt/data/kafka-logs/kafka-0] Starting up. (kafka.coordinator.log.LogCoordinator)
[2018-05-16 15:09:28,940] INFO [Log partition=test-0, dir=/opt/data/kafka-logs/kafka-0] Startup complete. (kafka.coordinator.log.LogCoordinator)
[2018-05-16 15:09:28,940] INFO [Log partition=test-0, dir=/opt/data/kafka-logs/kafka-0] [BrokerId=1] Removed 0 expired offsets. (kafka.coordinator.group.GroupMetadataManager)
[2018-05-16 15:09:28,940] INFO [Log partition=test-0, dir=/opt/data/kafka-logs/kafka-0] [BrokerId=1] Acquired new producerId block (producerId:3999) by writing to Zk with path version 5 (kafka.coordinator.transaction.ProducerIdManager)
[2018-05-16 15:09:28,940] INFO [Log partition=test-0, dir=/opt/data/kafka-logs/kafka-0] [BrokerId=1] Starting up. (kafka.coordinator.log.LogCoordinator)
[2018-05-16 15:09:28,940] INFO [Log partition=test-0, dir=/opt/data/kafka-logs/kafka-0] [BrokerId=1] Startup complete. (kafka.coordinator.log.LogCoordinator)
[2018-05-16 15:15:08,676] INFO [Transaction Marker Channel Manager 1]: Starting (kafka.coordinator.transaction.TransactionMarkerChannelManager)
[2018-05-16 15:15:09,267] INFO [/config/changes-event-process-thread]: Starting (kafka.common.ZkNodeChangeNotificationListener$ChangeEventProcessThread)
[2018-05-16 15:15:09,281] INFO Kafka version : 1.1.0 (org.apache.kafka.common.utils.AppInfoParser)
[2018-05-16 15:15:09,282] INFO Kafka commitId : fdcf75ea326b8e07 (org.apache.kafka.common.utils.AppInfoParser)
[2018-05-16 15:15:09,285] INFO [KafkaServer id=1] started (kafka.server.KafkaServer)
```

Create a replicated topic my-failsafe-topic

Now we will create a replicated topic that the console producers and console consumers can use.

Open a separate terminal and execute the following;

```
#/opt/kafka/bin/kafka-topics.sh --create --bootstrap-server kafka0:9092 --replication-factor 3 --partitions 13 --topic my-failsafe-topic
```

```
[root@tos ~]# /opt/kafka/bin/kafka-topics.sh --create --zookeeper localhost:2181 --replication-factor 3 --partitions 13 --topic my-failsafe-topic
Created topic "my-failsafe-topic".
[root@tos ~]#
```

Notice that the replication factor gets set to 3, and the topic name is **my-failsafe-topic**, and like before it has 13 partitions.

Start Kafka Consumer that uses Replicated Topic

Start the consumer with the script in a separate terminal;

```
#/opt/kafka/bin/kafka-console-consumer.sh --bootstrap-server kafka0:9094,kafka0:9092 --topic my-failsafe-topic --from-beginning
```

Notice that a list of Kafka servers is passed to **--bootstrap-server** parameter. Only, two of the three servers get passed that we ran earlier. Even though only one broker is needed, the

consumer client will learn about the other broker from just one server. Usually, you list multiple brokers in case there is an outage so that the client can connect.

```
[root@tos ~]# /opt/kafka/bin/kafka-console-consumer.sh --bootstrap-server localhost:9094,localhost:9092 --topic my-failsafe-topic
```

Start Kafka Producer that uses Replicated Topic

Next, we create a script that starts the producer. Then launch the producer with the script you create.

```
#/opt/kafka/bin/kafka-console-producer.sh --broker-list kafka0:9092,kafka0:9093 --topic my-failsafe-topic
```

Notice we start Kafka producer and pass it a list of Kafka Brokers to use via the parameter `--broker-list`.

Now use the start producer script to launch the producer as follows.

Now send messages

Now send some message from the producer to Kafka and see those messages consumed by the consumer.

Producer Console.

```
login as: root
root@10.10.20.24's password:
Last login: Wed May 16 15:20:40 2018 from 10.10.20.1
[root@tos ~]# /opt/kafka/bin/kafka-console-producer.sh --broker-list localhost:9092,localhost:9093 --topic my-failsafe-topic
>How are you?
>Great Kafka is working
>
```

Consumer Console

```
[root@tos ~]# /opt/kafka/bin/kafka-console-consumer.sh --bootstrap-server localhost:9094,localhost:9092 --topic my-failsafe-topic
How are you?
Great Kafka is working
```

Now Start two more consumers and send more messages.

Now Start two more consumers in their own terminal window and send more messages from the producer. (Replace the hostname of your server aka localhost)

Consumer 1.

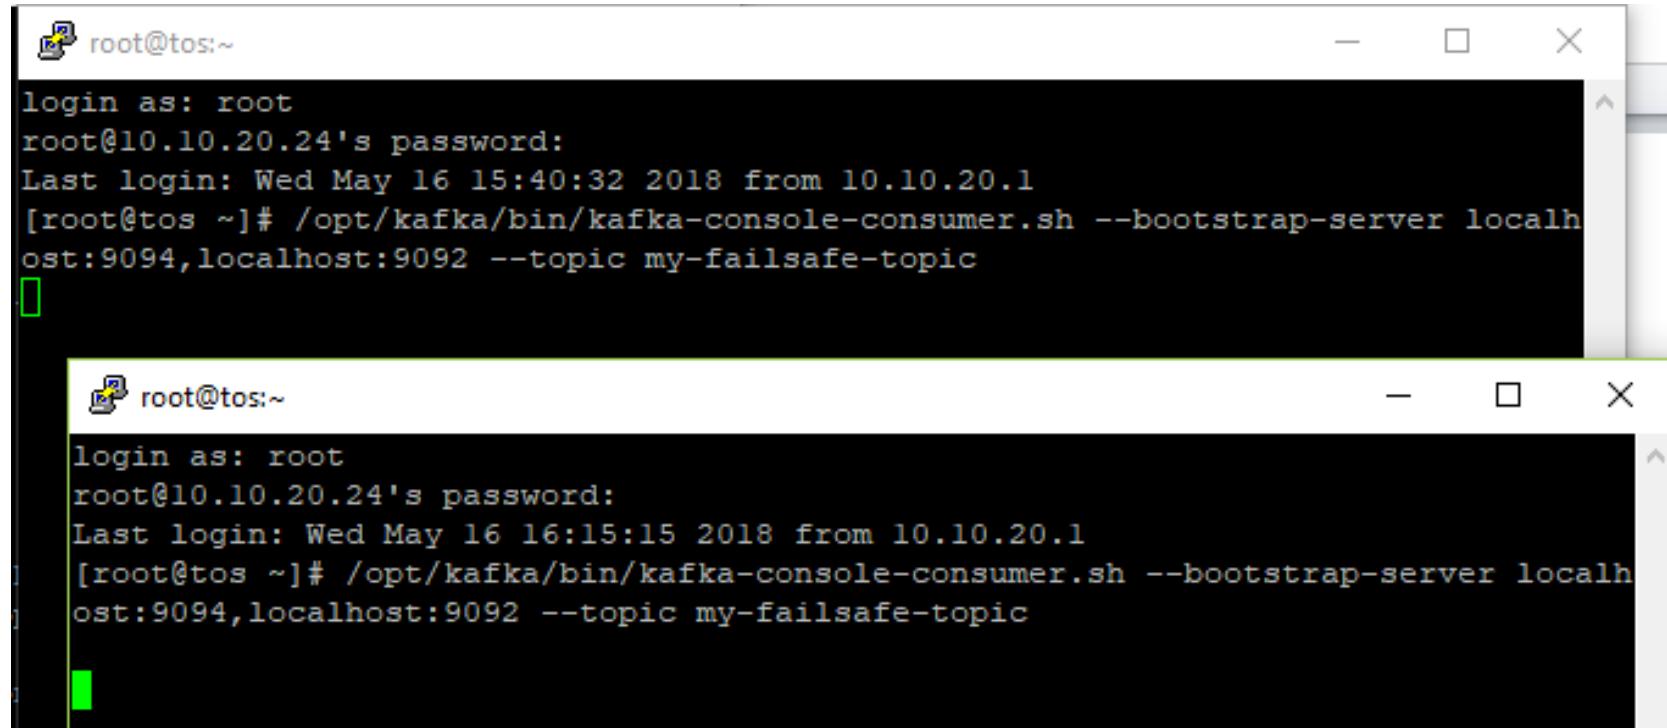
```
#/opt/kafka/bin/kafka-console-consumer.sh --bootstrap-server kafka0:9094,kafka0:9092
--topic my-failsafe-topic --from-beginning
```

Consumer 2.

```
#/opt/kafka/bin/kafka-console-consumer.sh --bootstrap-server kafka0:9094,kafka0:9092
--topic my-failsafe-topic --from-beginning
```

53 Kafka – Administration

The two consumer consoles will be as shown below



The image displays two separate terminal windows, each showing the command-line interface for a Kafka consumer. Both windows are titled 'root@tos:~' and show the same command being run:

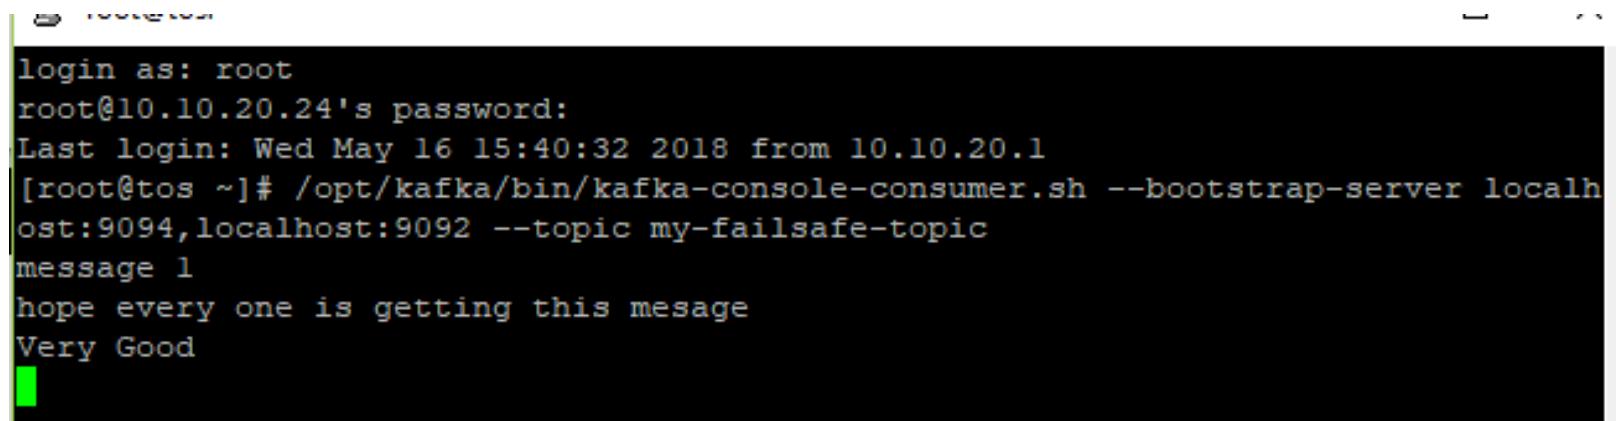
```
root@tos:~  
login as: root  
root@10.10.20.24's password:  
Last login: Wed May 16 15:40:32 2018 from 10.10.20.1  
[root@tos ~]# /opt/kafka/bin/kafka-console-consumer.sh --bootstrap-server localhost:9094,localhost:9092 --topic my-failsafe-topic
```

The windows have standard window controls (minimize, maximize, close) and are set against a light gray background.

Producer Console will be as shown below:

```
login as: root
root@10.10.20.24's password:
Last login: Wed May 16 15:20:40 2018 from 10.10.20.1
[root@tos ~]# /opt/kafka/bin/kafka-console-producer.sh --broker-list localhost:9092,localhost:9093 --topic my-failsafe-topic
>How are you?
>Great Kafka is working
>message 1
>hope every one is getting this mesage
>Very Good
>
```

Consumer Console 1st, you should be able to view the messages whatever you type on the producer console after the new consumer console was started.

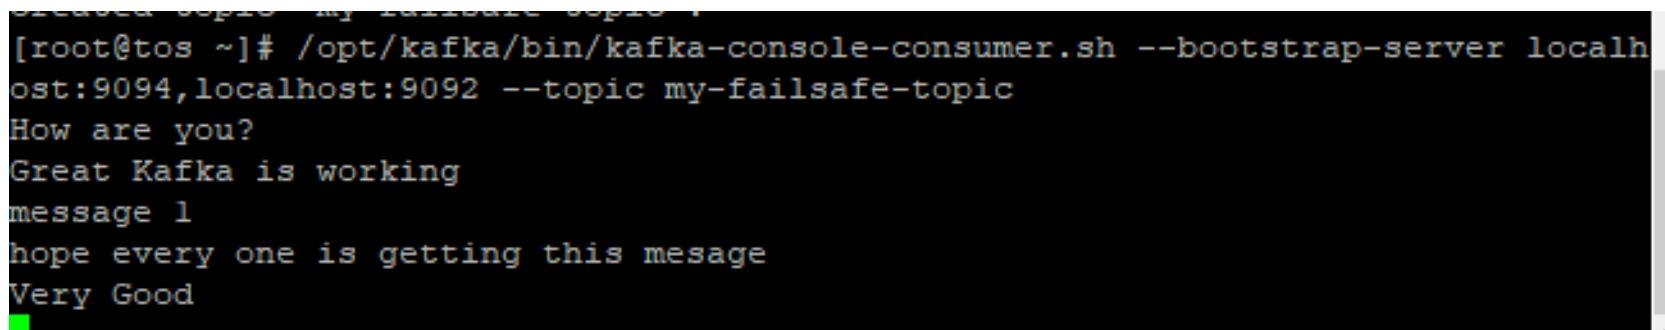


A terminal window showing the output of a Kafka consumer. The session starts with a root login, followed by the command to run the consumer. The consumer then prints several messages it has received from the topic 'my-failsafe-topic'.

```
root@tos ~]# /opt/kafka/bin/kafka-console-consumer.sh --bootstrap-server localhost:9094,localhost:9092 --topic my-failsafe-topic
message 1
hope every one is getting this mesage
Very Good
```

Consumer Console 2nd in new Terminal, similarly all the messages that were type on the producer console should be also display in the second console after it was started.

Consumer Console 2nd in new Terminal



A terminal window showing the output of a Kafka consumer. The session starts with a root login, followed by the command to run the consumer. The consumer prints several messages it has received from the topic 'my-failsafe-topic', which correspond to the messages sent from the first terminal.

```
[root@tos ~]# /opt/kafka/bin/kafka-console-consumer.sh --bootstrap-server localhost:9094,localhost:9092 --topic my-failsafe-topic
How are you?
Great Kafka is working
message 1
hope every one is getting this mesage
Very Good
```

Notice that the messages are sent to all of the consumers because each consumer is in a different consumer group.

Change consumer to be in their own consumer group.

Stop the producers and the consumers before, but leave Kafka and ZooKeeper running. You can use `ctrl + c`.

We want to put all of the consumers in same *consumer group*. This way the consumers will share the messages as each consumer in the *consumer group* will get its share of partitions.

Run the following scripts three times – from different console.

```
#/opt/kafka/bin/kafka-console-consumer.sh --bootstrap-server kafka0:9094,kafka0:9092  
--topic my-failsafe-topic --consumer-property group.id=mygroup
```

Notice that the script is the same as before except we added **--consumer-property group.id=mygroup** which will put every consumer that runs with this script into the **mygroup** consumer group.

Now we just run the producer and three consumers.

Run Producer Console

```
/opt/kafka/bin/kafka-console-producer.sh --broker-list kafka0:9092,kafka0:9093 --topic  
my-failsafe-topic
```

Now send eight messages from the Kafka producer console.

Producer Console

```
[root@tos ~]#  
[root@tos ~]#  
[root@tos ~]# /opt/kafka/bin/kafka-console-producer.sh --broker-list localhost:9  
092,localhost:9093 --topic my-failsafe-topic  
>message 1  
>message 2  
>message 3  
>message 4  
>message 5  
>message 6  
>message 7  
>message 8  
>
```

Notice that the messages are spread evenly among the consumers.

1st Kafka Consumer gets m3, m5

```
[root@tos ~]# /opt/kafka/bin/kafka-console-consumer.sh --bootstrap-server localhost:9094,localhost:9092 --topic my-failsafe-topic --consumer-property group.id=mygroup
message 5
message 6
message 7
```

Notice the first consumer gets messages m3 and m5.

2nd Kafka Consumer gets m2, m6

```
[root@tos ~]# /opt/kafka/bin/kafka-console-consumer.sh --bootstrap-server localhost:9094,localhost:9092 --topic my-failsafe-topic --consumer-property group.id=mygroup
message 2
message 4
message 8
```

Notice the second consumer gets messages m2 and m6.

3rd Kafka Consumer gets m1, m4, m7

```
[root@tos ~]# /opt/kafka/bin/kafka-console-consumer.sh --bootstrap-server localhost:9094,localhost:9092 --topic my-failsafe-topic --consumer-property group.id=mygroup
message 1
message 3
```

Notice the third consumer gets messages m1, m4 and m7.

Notice that each consumer in the group got a share of the messages.

Kafka Consumer Failover

Next, let's demonstrate consumer failover by killing one of the consumers and sending seven more messages. Kafka should divide up the work to the consumers that are running. First, kill the third consumer (CTRL-C in the consumer terminal does the trick).

```
[root@tos ~]# /opt/kafka/bin/kafka-console-consumer.sh --bootstrap-server localhost:9094,localhost:9092 --topic my-failsafe-topic --consumer-property group.id=mygroup
message 1
message 3
^CProcessed a total of 2 messages
[root@tos ~]#
```

Now send seven more messages from the Kafka console-producer.

Producer Console - send seven more messages m8 through m14

```
[root@tos ~]# /opt/kafka/bin/kafka-console-producer.sh --broker-list localhost:9092,localhost:9093 --topic my-failsafe-topic
>message 1
>message 2
>message 3
>message 4
>message 5
>message 6
>message 7
>message 8
>
>message 9
>m 10
>m11
>m 12
>m 13
>m 14
>m 15
>
```

Notice that the messages are spread evenly among the remaining consumers.

1st Kafka Consumer gets m8, m9, m11, m14

```
[root@tos ~]# /opt/kafka/bin/kafka-console-consumer.sh --bootstrap-server local
host:9094,localhost:9092 --topic my-failsafe-topic --consumer-property group.
id=mygroup
message 2
message 4
message 8

message 9
m 12
m 13
m 14
```

2nd Kafka Consumer gets m10, m12, m13

```
[root@tos ~]# /opt/kafka/bin/kafka-console-consumer.sh --bootstrap-server local
host:9094,localhost:9092 --topic my-failsafe-topic --consumer-property group.
id=mygroup
message 5
message 6
message 7
m 10
m11
m 15
```

We killed one consumer, sent seven more messages, and saw Kafka spread the load to remaining consumers. ***Kafka consumer failover works!***

Create Kafka Describe Topic Script

You can use `kafka-topics.sh` to see how the Kafka topic is laid out among the Kafka brokers. The `--describe` will show partitions, ISRs, and broker partition leadership.

```
#/opt/kafka/bin/kafka-topics.sh --describe --topic my-failsafe-topic --bootstrap-server kafka:9092
```

```
[root@tos ~]# /opt/kafka/bin/kafka-topics.sh --describe --topic my-failsafe-topic --zookeeper localhost:2181
Topic:my-failsafe-topic PartitionCount:13      ReplicationFactor:3      Configs:
      Topic: my-failsafe-topic      Partition: 0      Leader: 2      Replicas: 2,0,1 Isr: 0,1,2
      Topic: my-failsafe-topic      Partition: 1      Leader: 0      Replicas: 0,1,2 Isr: 0,1,2
      Topic: my-failsafe-topic      Partition: 2      Leader: 1      Replicas: 1,2,0 Isr: 0,1,2
      Topic: my-failsafe-topic      Partition: 3      Leader: 2      Replicas: 2,1,0 Isr: 0,1,2
      Topic: my-failsafe-topic      Partition: 4      Leader: 0      Replicas: 0,2,1 Isr: 0,1,2
      Topic: my-failsafe-topic      Partition: 5      Leader: 1      Replicas: 1,0,2 Isr: 0,1,2
      Topic: my-failsafe-topic      Partition: 6      Leader: 2      Replicas: 2,0,1 Isr: 0,1,2
      Topic: my-failsafe-topic      Partition: 7      Leader: 0      Replicas: 0,1,2 Isr: 0,1,2
      Topic: my-failsafe-topic      Partition: 8      Leader: 1      Replicas: 1,2,0 Isr: 0,1,2
      Topic: my-failsafe-topic      Partition: 9      Leader: 2      Replicas: 2,1,0 Isr: 0,1,2
      Topic: my-failsafe-topic      Partition: 10     Leader: 0      Replicas: 0,2,1 Isr: 0,1,2
      Topic: my-failsafe-topic      Partition: 11     Leader: 1      Replicas: 1,0,2 Isr: 0,1,2
      Topic: my-failsafe-topic      Partition: 12     Leader: 2      Replicas: 2,0,1 Isr: 0,1,2
[root@tos ~]#
```

Run describe-topics

We are going to lists which broker owns (leader of) which partition, and list replicas and ISRs of each partition. ISRs are replicas that are up to date. Remember there are 13 topics.

Topology of Kafka Topic Partition Ownership

Notice how each broker gets a share of the partitions as leaders and followers. Also, see how Kafka replicates the partitions on each broker.

Test Broker Failover by killing 1st server

Let's kill the first broker, and then test the failover.

Kill the first broker

```
$ kill `ps aux | grep java | grep server-o.properties | tr -s " " | cut -d " " -f2`
```

```
[2018-05-17 17:38:07,157] INFO [ThrottledRequestReaper-Request]: Shutdown completed (kafka.server.ClientQuotaManager$ThrottledRequestReaper)
[2018-05-17 17:38:07,158] INFO [SocketServer brokerId=0] Shutting down socket server (kafka.network.SocketServer)
[2018-05-17 17:38:07,267] INFO [SocketServer brokerId=0] Shutdown completed (kafka.network.SocketServer)
[2018-05-17 17:38:07,286] INFO [KafkaServer id=0] shut down completed (kafka.server.KafkaServer)
[root@tos kafka-config]#
```

You can stop the first broker by hitting CTRL-C in the broker terminal or by running the above command.

Now that the first Kafka broker has stopped, let's use Kafka [topics describe](#) to see that new leaders were elected!

Run `describe-topics` again to see leadership change

```
[root@tos ~]# /opt/kafka/bin/kafka-topics.sh --describe --topic my-failsafe-topic --zookeeper localhost:2181
Topic:my-failsafe-topic PartitionCount:13      ReplicationFactor:3      Configs:
  Topic: my-failsafe-topic      Partition: 0      Leader: 2      Replicas: 2,0,1 Isr: 1,2
  Topic: my-failsafe-topic      Partition: 1      Leader: 1      Replicas: 0,1,2 Isr: 1,2
  Topic: my-failsafe-topic      Partition: 2      Leader: 1      Replicas: 1,2,0 Isr: 1,2
  Topic: my-failsafe-topic      Partition: 3      Leader: 2      Replicas: 2,1,0 Isr: 1,2
  Topic: my-failsafe-topic      Partition: 4      Leader: 2      Replicas: 0,2,1 Isr: 1,2
  Topic: my-failsafe-topic      Partition: 5      Leader: 1      Replicas: 1,0,2 Isr: 1,2
  Topic: my-failsafe-topic      Partition: 6      Leader: 2      Replicas: 2,0,1 Isr: 1,2
  Topic: my-failsafe-topic      Partition: 7      Leader: 1      Replicas: 0,1,2 Isr: 1,2
  Topic: my-failsafe-topic      Partition: 8      Leader: 1      Replicas: 1,2,0 Isr: 1,2
  Topic: my-failsafe-topic      Partition: 9      Leader: 2      Replicas: 2,1,0 Isr: 1,2
  Topic: my-failsafe-topic      Partition: 10     Leader: 2      Replicas: 0,2,1 Isr: 1,2
  Topic: my-failsafe-topic      Partition: 11     Leader: 1      Replicas: 1,0,2 Isr: 1,2
  Topic: my-failsafe-topic      Partition: 12     Leader: 2      Replicas: 2,0,1 Isr: 1,2
[root@tos ~]#
```

Notice how Kafka spreads the leadership over the 2nd and 3rd Kafka brokers.

Create consolidated scripts as mention below. It will help you to start zookeeper and kafka using a single script.

Start 3 Brokers.

```
#vi /opt/scripts/start3Brokers.sh
#!/usr/bin/env bash
# Start Zookeeper.
/opt/zookeeper/bin/zkServer.sh start
```

#Start Kafka Server 0

```
/opt/kafka/bin/kafka-server-start.sh "/opt/kafka-config/config/server-0.properties" &
```

#Start Kafka Server 1

```
/opt/kafka/bin/kafka-server-start.sh "/opt/kafka-config/config/server-1.properties" &
```

```
#Start Kafka Server 2  
/opt/kafka/bin/kafka-server-start.sh "/opt/kafka-config/config/server-2.properties" &  
echo "Start zookeeper & 3 Brokers Successfully"
```

to Stop All 3 Brokers.

```
# vi stop3Brokers.sh  
#!/usr/bin/env bash  
# Stop Zookeeper & 3 Servers.  
/opt/zookeeper/bin/zkServer.sh stop  
/opt/kafka/bin/kafka-server-stop.sh  
echo "Stop zookeeper & 3 servers Successfully"
```

----- Lab Ends here -----

8. Securing Kafka ACL – 60 Minutes

Kafka SASL/PLAIN without SSL

We will perform the following activity in a single node broker only. We will enable ACL in a single broker.

To run a secure broker, two steps need to be performed.

1. Configure the Kafka brokers with ACL using JAAS file
2. Kafka Clients – Pass the Credentials

First, we need to let the broker know authorized users' credentials. This will be stored in a JAAS file.

1. Configure the Kafka brokers and Kafka Clients

Add a JAAS configuration file for each Kafka broker. Create a kafka_plain_jaas.conf file as specified below:

Prepare the /opt/kafka/config/kafka_jaas.conf with below contents. Use vi editor.

```
KafkaServer {  
    org.apache.kafka.common.security.plain.PlainLoginModule required  
        username="admin"  
        password="admin"
```

```
user_admin="admin"  
user_alice="alice"  
user_bob="bob"  
user_charlie="charlie";  
};
```

Let's understand the content of ***kafka_plain_jaas.conf*** file and how Kafka Brokers and Kafka Clients use it.

KafkaServer Section:

The KafkaServer section defines four users: admin, alice, bob and charlie. The properties username and password are used by the broker to initiate connections to other brokers. In this example, admin is the user for inter-broker communication. The set of properties user_{userName} defines the passwords for all users that connect to the broker and the broker validates all client connections including those from other brokers using these properties.

This file needs to be passed in as a JVM config option when running the broker, using -Djava.security.auth.login.config=[path_to_jaas_file].

Create an initialize scripts in /opt/scripts/kafkasecurity.sh file with below contents to setup kafka_jaas.conf file location

```
#vi /opt/scripts/kafkasecurity.sh
```

Update with the following content.

```
#!/usr/bash
export KAFKA_HOME=/opt/kafka
export ZOOKEEPER_HOME=/opt/zookeeper
export KAFKA_PLAIN_PARAMS="-Djava.security.auth.login.config=/opt/kafka/config/kafka_jaas.conf"
export KAFKA_OPTS="$KAFKA_PLAIN_PARAMS $KAFKA_OPTS"
export PATH="$PATH:$ZOOKEEPER_HOME/bin:$KAFKA_HOME/bin:$JAVA_HOME/bin"
```



```
...ata/kafka-logs (com.docker.cli)  ●  961 ...pt/kafka/config (com.docker.cli)  ●  962 ...0:/opt/scripts/cer (com.docker.cli)  963
#!/usr/bash
export KAFKA_HOME=/opt/kafka
export ZOOKEEPER_HOME=/opt/zookeeper
export KAFKA_PLAIN_PARAMS="-Djava.security.auth.login.config=/opt/kafka/config/kafka_jaas.conf"
export KAFKA_OPTS="$KAFKA_PLAIN_PARAMS $KAFKA_OPTS"
export PATH="$PATH:$ZOOKEEPER_HOME/bin:$KAFKA_HOME/bin:$JAVA_HOME/bin"
```

Use the following interpreter in case of **bash** interpreter issue.

```
#!/bin/bash
```

Second, the following needs to be added to the broker properties file (e.g. server.properties) it defines the accepted protocol and also the ACL authorizer used by the broker:

Step 4: Update **/opt/kafka/config/server.properties** for all brokers. In our case only one broker.

```
# Use below for SASL/PLAIN only (No SSL)
# Using SASL_PLAINTEXT as we do not have SSL
authorizer.class.name=kafka.security.authorizer.AclAuthorizer

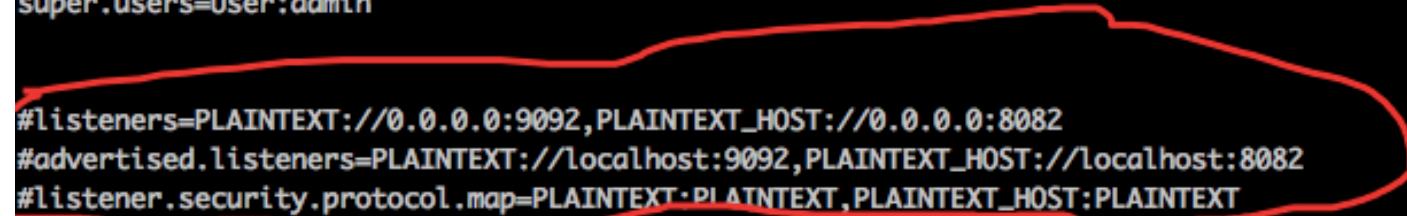
listeners=SASL_PLAINTEXT://kafka0:9092
security.inter.broker.protocol=SASL_PLAINTEXT
sasl.mechanism.inter.broker.protocol=PLAIN
sasl.enabled.mechanisms=PLAIN
advertised.listeners=SASL_PLAINTEXT://kafka0:9092
super.users=User:admin
```

```
#      listeners = PLAINTEXT://your.host.name:9092
authorizer.class.name=kafka.security.auth.SimpleAclAuthorizer
listeners=SASL_PLAINTEXT://tos.master.com:9092
security.inter.broker.protocol= SASL_PLAINTEXT
sasl.mechanism.inter.broker.protocol=PLAIN
sasl.enabled.mechanisms=PLAIN
# Hostname and port the broker will advertise to producers and consumers. If not
# set,
# it uses the value for "listeners" if configured. Otherwise, it will use the v
alue
# returned from java.net.InetAddress.getCanonicalHostName().
advertised.listeners=SASL_PLAINTEXT://tos.master.com:9092

super.users=User:admin
# Map listener names to security protocols - the default is for them to be the s
```

Comment the following line.

```
# use Ssl for security layer only (no SSL)
# Using SASL_PLAINTEXT as we do not have SSL
#authorizer.class.name=kafka.security.auth.SimpleAclAuthorizer
authorizer.class.name=kafka.security.authorizer.AclAuthorizer
listeners=SASL_PLAINTEXT://localhost:9092
security.inter.broker.protocol=SASL_PLAINTEXT
sasl.mechanism.inter.broker.protocol=PLAIN
sasl.enabled.mechanisms=PLAIN
advertised.listeners=SASL_PLAINTEXT://localhost:9092
super.users=User:admin
```



```
#listeners=PLAINTEXT://0.0.0.0:9092,PLAINTEXT_HOST://0.0.0.0:8082
#advertised.listeners=PLAINTEXT://localhost:9092,PLAINTEXT_HOST://localhost:8082
#listener.security.protocol.map=PLAINTEXT:PLAINTEXT,PLAINTEXT_HOST:PLAINTEXT
```

This is to initialize the environment HOME.

```
#cd /opt/scripts  
#source ./kafkasecurity.sh
```

Before starting the server, let us clean the zookeeper/kafka log data so that there won't be any conflict with configuration of the previous cluster (**/opt/data/zookeeper** and **(/opt/data/kafka-logs)**).

```
[root@tos zookeeper]# ls  
version-2  zookeeper_server.pid  
[root@tos zookeeper]# pwd  
/opt/data/zookeeper  
[root@tos zookeeper]# ls  
version-2  zookeeper_server.pid  
[root@tos zookeeper]# rm -fr *  
[root@tos zookeeper]# █
```

Delete the data folder of the zookeeper as shown above

Restart ZK/Kafka Services or

Start our broker (from the console that executed : source ./ kafkasecurity.sh)

```
# sh startABroker.sh
```

```
[root@tos scripts]# sh startABroker.sh
ZooKeeper JMX enabled by default
Using config: /opt/zookeeper/bin/..../conf/zoo.cfg
Starting zookeeper ... STARTED
Started Successfully
[root@tos scripts]# jps
2551 Kafka
2584 Jps
2286 QuorumPeerMain
[root@tos scripts]#
```

Validate the logs: **/opt/kafka/logs/server.log** in case there is any error.

Note : You can include sh startABroker.sh inside the kafkaSecurity file in case the environments are not set properly in some OS.

```
(kafka.server.KafkaConfig)
[2018-06-18 15:33:12,246] INFO KafkaConfig values:
    advertised.host.name = null
    advertised.listeners = SASL_PLAINTEXT://tos.master.com:9092
    advertised.port = null
    alter.config.policy.class.name = null
    alter.log.dirs.replication.quota.window.num = 11
    alter.log.dirs.replication.quota.window.size.seconds = 1
    authorizer.class.name =
    auto.create.topics.enable = true
    auto.leader.rebalance.enable = true
    background.threads = 10
    broker.id = 1
    broker.id.generation.enable = true
    broker.rack = null
    compression.type = producer
    connections.max.idle.ms = 600000
    controlled.shutdown.enable = true
    controlled.shutdown.max.retries = 3
    controlled.shutdown.retry.backoff.ms = 5000
```

All the commands shown below should be executed from the path `/opt/scripts`.

Create a topic:

```
#/opt/kafka/bin/kafka-topics.sh --create --bootstrap-server kafka0:9092 --replication-factor 1 --partitions 1 --topic plain-topic
```

```
[root@kafka0 kafka-logs]# /opt/kafka/bin/kafka-topics.sh --create --bootstrap-server localhost:9092 --replication-factor 1 --partitions 1 --topic plain-topic
Error while executing topic command : Timed out waiting for a node assignment. Call: createTopics
[2022-05-17 16:42:55,686] ERROR org.apache.kafka.common.errors.TimeoutException: Timed out waiting for a
node assignment. Call: createTopics
(kafka.admin.TopicCommand$)
[root@kafka0 kafka-logs]#
```

It should generate an error like as shown above.

Run Kafka console producer.

Before running the Kafka console Producer configure the producer.properties file as shown:

vi producer.properties

Update with the following entries.

```
sasl.jaas.config=org.apache.kafka.common.security.plain.PlainLoginModule required \
username="alice" \
password="alice";
security.protocol=SASL_PLAINTEXT
sasl.mechanism=PLAIN
```

----- Code Ends Here // Exclude this line.

```
#/opt/kafka/bin/kafka-console-producer.sh --broker-list kafka:9092 --topic plain-topic \
--producer.config producer.properties
```

Send some messages in the producer console.

Message 1

Message 2

Message 3

```
[2020-08-30 19:14:56,662] WARN [Producer clientId=console-producer] Error while fetching metadata with correlation id 3 : {plain-topic=TOPIC_AUTHORIZATION_FAILED} (org.apache.kafka.clients.NetworkClient)
[2020-08-30 19:14:56,668] ERROR [Producer clientId=console-producer] Topic authorization failed for topics [plain-topic] (org.apache.kafka.clients.Metadata)
[2020-08-30 19:14:56,670] ERROR Error when sending message to topic plain-topic with key: null, value: 9 bytes with error: (org.apache.kafka.clients.producer.internals.ErrorLoggingCallback)
org.apache.kafka.common.errors.TopicAuthorizationException: Not authorized to access topics: [plain-topic]
>[2020-08-30 19:14:56,768] WARN [Producer clientId=console-producer] Error while fetching metadata with correlation id 4 : {plain-topic=TOPIC_AUTHORIZATION_FAILED} (org.apache.kafka.clients.NetworkClient)
[2020-08-30 19:14:56,769] ERROR [Producer clientId=console-producer] Topic authorization failed for topics [plain-topic] (org.apache.kafka.clients.Metadata)
[2020-08-30 19:14:56,769] ERROR Error when sending message to topic plain-topic with key: null, value: 9 bytes with error: (org.apache.kafka.clients.producer.internals.ErrorLoggingCallback)
org.apache.kafka.common.errors.TopicAuthorizationException: Not authorized to access topics: [plain-topic]
>[2020-08-30 19:14:56,872] WARN [Producer clientId=console-producer] Error while fetching metadata with correlation id 5 : {plain-topic=TOPIC_AUTHORIZATION_FAILED} (org.apache.kafka.clients.NetworkClient)
[2020-08-30 19:14:56,873] ERROR [Producer clientId=console-producer] Topic authorization failed for topics [plain-topic] (org.apache.kafka.clients.Metadata)
[2020-08-30 19:14:56,873] ERROR Error when sending message to topic plain-topic with key: null, value: 9 bytes with error: (org.apache.kafka.clients.producer.internals.ErrorLoggingCallback)
org.apache.kafka.common.errors.TopicAuthorizationException: Not authorized to access topics: [plain-topic]
>
```

Error since we have not authorized to access the topic.

The security configuration still does not give specific permissions to our Kafka users (except for admin who is a super user). These permissions are defined using the ACL command (bin/kafka-acls.sh). To verify existing ACLs run:

```
#cd /opt/scripts  
#vi admin.properties
```

Update the following content in it.
// Begin - Exclude it

```
sasl.jaas.config=org.apache.kafka.common.security.plain.PlainLoginModule required \  
username="admin" \  
password="admin";  
security.protocol=SASL_PLAINTEXT  
sasl.mechanism=PLAIN
```

// Ends Here – Exclude it

Verify the ACL configure using the Admin credentials.

```
#/opt/kafka/bin/kafka-acls.sh --bootstrap-server kafka:9092 --command-config  
admin.properties --list
```

This returns no ACL definitions. You have handled authentication, but have not provided any authorization rules to define the users able run specific APIs and access certain Kafka resources.

```
[root@kafka0 kafka-logs]# cd /opt/scripts/
[root@kafka0 scripts]# vi admin.properties
[root@kafka0 scripts]# /opt/kafka/bin/kafka-acls.sh --bootstrap-server localhost:9092 --command-config admin.properties --list
[root@kafka0 scripts]# 
```

Assign the permission:

```
# /opt/kafka/bin/kafka-acls.sh --bootstrap-server kafka0:9092 --command-config
admin.properties --add --allow-principal User:alice --operation All --topic plain-topic
```

```
[root@kafka0 config]# cd /opt/scripts
[root@kafka0 scripts]# /opt/kafka/bin/kafka-acls.sh --bootstrap-server localhost:9092 --command-config admin.properties --
add --allow-principal User:alice --operation All --topic plain-topic
Adding ACLs for resource `ResourcePattern(resourceType=TOPIC, name=plain-topic, patternType=LITERAL)`:
  (principal=User:alice, host=*, operation=ALL, permissionType=ALLOW)

Current ACLs for resource `ResourcePattern(resourceType=TOPIC, name=plain-topic, patternType=LITERAL)`:
  (principal=User:alice, host=*, operation=WWRITE, permissionType=ALLOW)
  (principal=User:alice, host=*, operation=ALL, permissionType=ALLOW)

[root@kafka0 scripts]# 
```

Now you can verify the acl granted to alice user

```
# /opt/kafka/bin/kafka-acls.sh --bootstrap-server kafka0:9092 --command-config admin.properties --list
```

```
[root@kafka0 scripts]# /opt/kafka/bin/kafka-acls.sh --bootstrap-server localhost:9092 --command-config admin.properties --list
Current ACLs for resource `ResourcePattern(resourceType=TOPIC, name=plain-topic, patternType=LITERAL)`:
  (principal=User:alice, host=*, operation=WRITE, permissionType=ALLOW)
  (principal=User:alice, host=*, operation=ALL, permissionType=ALLOW)

[root@kafka0 scripts]#
```

Now you can try sending message again from the producer console. Type one or two messages in the producer console.

```
# cd /opt/scripts
#/opt/kafka/bin/kafka-console-producer.sh --broker-list kafka0:9092 --topic plain-topic --producer.config producer.properties
```

Send some messages in the producer console.

Hello Message

```
[root@kafka0 scripts]# /opt/kafka/bin/kafka-console-producer.sh --broker-list localhost:9092 --topic plain-topic --producer
.config producer.properties
>Hello
>Message
>
```

Now, it should be able to sent message to the topic as shown above.

Now we can try consuming these messages further from the consumer console, open a separate terminal for this.

Run Kafka console consumer

Consuming from Topic using bob user, hence grant ACL to bob.

Next, you need to let user bob consume (or fetch) from topic plain-topic using the Fetch API, as a member of the bob-group consumer group. Bob's ACL for fetching from topic test is:

Principal bob is Allowed Operation Read From Host * On Topic plain-topic .

or

```
# /opt/kafka/bin/kafka-acls.sh --bootstrap-server=kafka0:9092 -command-config admin.properties --add --allow-principal User:bob --operation Read --topic plain-topic
```

```
[root@kafka0 scripts]# /opt/kafka/bin/kafka-acls.sh --bootstrap-server=localhost:9092 -command-config admin.properties --add --allow-principal User:bob --operation Read --topic plain-topic
Adding ACLs for resource `ResourcePattern(resourceType=TOPIC, name=plain-topic, patternType=LITERAL)`:
  (principal=User:bob, host=*, operation=READ, permissionType=ALLOW)

Current ACLs for resource `ResourcePattern(resourceType=TOPIC, name=plain-topic, patternType=LITERAL)`:
  (principal=User:bob, host=*, operation=READ, permissionType=ALLOW)
  (principal=User:alice, host=*, operation=WRITE, permissionType=ALLOW)
  (principal=User:alice, host=*, operation=ALL, permissionType=ALLOW)

[root@kafka0 scripts]#
```

Bob needs a second ACL for committing offsets to group bob-group (using the OffsetCommit API):

Principal bob is Allowed Operation Read From Host * On Group bob-group.

or

```
$ /opt/kafka/bin/kafka-acls.sh --bootstrap-server=kafka0:9092 -command-config admin.properties --add --allow-principal User:bob --operation Read --group bob-group
```

```
[root@kafka0 scripts]# /opt/kafka/bin/kafka-acls.sh --bootstrap-server=localhost:9092 -command-config admin.properties --add --allow-principal User:bob --operation Read --group bob-group
Adding ACLs for resource `ResourcePattern(resourceType=GROUP, name=bob-group, patternType=LITERAL)`:
  (principal=User:bob, host=*, operation=READ, permissionType=ALLOW)

Current ACLs for resource `ResourcePattern(resourceType=GROUP, name=bob-group, patternType=LITERAL)`:
  (principal=User:bob, host=*, operation=READ, permissionType=ALLOW)

[root@kafka0 scripts]#
```

By granting these permissions to Bob, he can now consume messages from topic test as a member of bob-group.

Before running Kafka console consumer configure the consumer.properties file as shown:

```
# vi consumer.properties
```

```
sasl.jaas.config=org.apache.kafka.common.security.plain.PlainLoginModule required \
  username="bob" \
  password="bob";
security.protocol=SASL_PLAINTEXT
sasl.mechanism=PLAIN
```

```
$ /opt/kafka/bin/kafka-console-consumer.sh --bootstrap-server kafka:9092 --topic plain-topic --group bob-group --from-beginning --consumer.config consumer.properties
```

```
[root@kafka0 scripts]# /opt/kafka/bin/kafka-console-consumer.sh --bootstrap-server localhost:9092 --topic plain-topic --group bob-group --from-beginning --consumer.config consumer.properties
hi
Hello
Message
```

Sent a few messages from the producer console, and you will see that messages are being consumed on the console...

Describing consumer groups

Open a new console for the Charlie user.

Lastly, user charlie needs permission to retrieve committed offsets from group bob-group (using the OffsetFetch API). According to the table above, Charlie's first ACL for fetching offsets from this consumer group is:

Principal charlie is Allowed Operation Describe From Host * On Group bob-group.

or

```
$ /opt/kafka/bin/kafka-acls.sh --bootstrap-server=kafkao:9092 -command-config admin.properties --add --allow-principal User:charlie --operation Describe --group bob-group
```

```
[root@kafka0 scripts]# /opt/kafka/bin/kafka-acls.sh --bootstrap-server=localhost:9092 -command-config admin.properties --add --allow-principal User:charlie --operation Describe --group bob-group
Adding ACLs for resource `ResourcePattern(resourceType=GROUP, name=bob-group, patternType=LITERAL)`:
  (principal=User:charlie, host=*, operation=DESCRIBE, permissionType=ALLOW)

Current ACLs for resource `ResourcePattern(resourceType=GROUP, name=bob-group, patternType=LITERAL)`:
  (principal=User:charlie, host=*, operation=DESCRIBE, permissionType=ALLOW)
  (principal=User:bob, host=*, operation=READ, permissionType=ALLOW)

[root@kafka0 scripts]#
```

Principal charlie is Allowed Operation Describe From Host * On Topic test.

or

```
$ /opt/kafka/bin/kafka-acls.sh --bootstrap-server=kafkao:9092 -command-config admin.properties --add --allow-principal User:charlie --operation Describe --topic plain-topic
```

```
[root@kafka0 scripts]# /opt/kafka/bin/kafka-acls.sh --bootstrap-server=localhost:9092 --command-config admin.properties --add --allow-principal User:charlie --operation Describe --topic plain-topic
Adding ACLs for resource `ResourcePattern(resourceType=TOPIC, name=plain-topic, patternType=LITERAL)`:
(principal=User:charlie, host=*, operation=DESCRIBE, permissionType=ALLOW)

Current ACLs for resource `ResourcePattern(resourceType=TOPIC, name=plain-topic, patternType=LITERAL)`:
(principal=User:bob, host=*, operation=READ, permissionType=ALLOW)
(principal=User:charlie, host=*, operation=DESCRIBE, permissionType=ALLOW)
(principal=User:alice, host=*, operation=WRITE, permissionType=ALLOW)
(principal=User:alice, host=*, operation=ALL, permissionType=ALLOW)

[root@kafka0 scripts]#
```

The following file contains the client configuration for group.

```
# vi group.properties
```

```
sasl.jaas.config=org.apache.kafka.common.security.plain.PlainLoginModule required \
username="charlie" \
password="charlie";
security.protocol=SASL_PLAINTEXT
sasl.mechanism=PLAIN
```

Now Charlie is able to get the proper listing of offsets in the group:

```
$ /opt/kafka/bin/kafka-consumer-groups.sh --bootstrap-server kafkao:9092 --describe --group bob-group --command-config group.properties
```

```
[root@tos scripts]# vi group.properties
[root@tos scripts]# /opt/kafka/bin/kafka-consumer-groups.sh --bootstrap-server tos.master.
com:9092 --describe --group bob-group --command-config group.properties
Note: This will not show information about old Zookeeper-based consumers.

TOPIC      PARTITION  CURRENT-OFFSET  LOG-END-OFFSET  LAG          CONSUMER-ID
                  HOST        CLIENT-ID
plain-topic    0          5              5            0  consumer-l-6450
d497-85b1-4113-80bc-985c4045b185 /10.10.20.24  consumer-l
[root@tos scripts]#
```

Or

```
[root@kafka0 scripts]# vi group.properties
[root@kafka0 scripts]# /opt/kafka/bin/kafka-consumer-groups.sh --bootstrap-server localhost:9092 --describe --group bob-grou
p --command-config group.properties

Consumer group 'bob-group' has no active members.

GROUP      TOPIC      PARTITION  CURRENT-OFFSET  LOG-END-OFFSET  LAG          CONSUMER-ID  HOST
CLIENT-ID
bob-group  plain-topic  0          4              4            0          -          -          -
[root@kafka0 scripts]#
```

As you can see above the details of the consumer group, Here the log end offset is showing as 3/5 that is the current offset are same with that of the end offset. This is an ideal cluster, where there is no lagging in the production and consumption.

The above ACLs grants enough permissions for this use case to run.
To summarize, configured ACLs execute the following command.

```
$ /opt/kafka/bin/kafka-acls.sh --bootstrap-server=kafkao:9092 --command-config admin.properties --list
```

```
[root@kafka0 scripts]# /opt/kafka/bin/kafka-acls.sh --bootstrap-server=localhost:9092 --command-config admin.properties --list
Current ACLs for resource `ResourcePattern(resourceType=GROUP, name=bob-group, patternType=LITERAL)`:
  (principal=User:charlie, host=*, operation=DESCRIBE, permissionType=ALLOW)
  (principal=User:bob, host=*, operation=READ, permissionType=ALLOW)

Current ACLs for resource `ResourcePattern(resourceType=TOPIC, name=plain-topic, patternType=LITERAL)`:
  (principal=User:bob, host=*, operation=READ, permissionType=ALLOW)
  (principal=User:charlie, host=*, operation=DESCRIBE, permissionType=ALLOW)
  (principal=User:alice, host=*, operation=WRITE, permissionType=ALLOW)
  (principal=User:alice, host=*, operation=ALL, permissionType=ALLOW)

[root@kafka0 scripts]#
```

Errors:

```
[root@tos scripts]# /opt/kafka/bin/kafka-console-consumer.sh --bootstrap-server tos.master.com:9092 --topic plain-topic --from-beginning --consumer.config consumer.properties
[2018-06-27 21:37:46,223] WARN [Consumer clientId=consumer-1, groupId=console-consumer-47559] Error while fetching metadata with correlation id 2 : {plain-topic=TOPIC_AUTHORIZATION_FAILED} (org.apache.kafka.clients.NetworkClient)
[2018-06-27 21:37:46,232] ERROR Unknown error when running consumer: (kafka.tools.ConsoleConsumer$)
org.apache.kafka.common.errors.GroupAuthorizationException: Not authorized to access group: console-consumer-47559
```

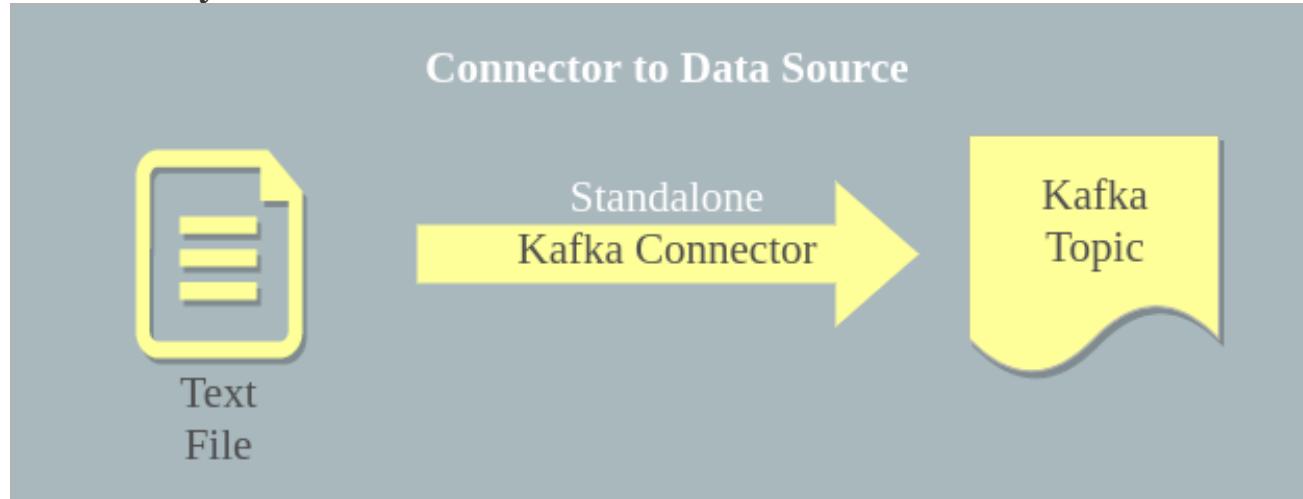
Grant proper authorization.

Ensure that kafkasecurity.sh is only executed for the server start up scripts not in any of the other console.

Lab Ends Here

9. Kafka Connector (File) - 60 Minutes

Apache Kafka Connector – Connectors are the components of Kafka that could be setup to listen the changes that happen to a data source like a file or database, and pull in those changes automatically.



When working with Kafka you might need to write data from a local file to a Kafka topic. This is actually very easy to do with Kafka Connect. Kafka Connect is a framework that provides scalable and reliable streaming of data to and from Apache Kafka. With Kafka Connect, writing a file's content to a topic requires only a few simple steps

Starting Kafka and Zookeeper

We will use the standalone Broker this example.

```
#cd /opt/scripts  
#sh startABroker.sh
```

```
[root@tos scripts]# ls
custom-reassignment.json  server2.log      stopZookeeper.sh
old.json                  start3Brokers.sh topics-to-move.json
server0.log                startABroker.sh  zookeeper.out
server1.log                stop3Brokers.sh
[root@tos scripts]# sh startABroker.sh
ZooKeeper JMX enabled by default
Using config: /opt/zookeeper/bin/..../conf/zoo.cfg
Starting zookeeper ... STARTED
Started Successfully
[root@tos scripts]# jps
2209 QuorumPeerMain
2474 Kafka
2493 Jps
[root@tos scripts]#
```

Creating a Topic to Write to, the message will be fetch from the file and publish to the topic.

```
#cd /opt/scripts
#/opt/kafka/bin/kafka-topics.sh \
--create \
--bootstrap-server kafka:9092 \
--replication-factor 1 \
--partitions 1 \
--topic my-kafka-connect
```

```
[root@tos bin]#
[root@tos bin]# cd /opt/scripts
[root@tos scripts]# /opt/kafka/bin/kafka-topics.sh \
>   --create \
>   --zookeeper tos.master.com:2181 \
>   --replication-factor 1 \
>   --partitions 1 \
>   --topic my-kafka-connect
Created topic "my-kafka-connect".
[root@tos scripts]#
```

Creating a Source Config File

Since we are reading the contents of a local file and writing to Kafka, this file is considered our “source”. Therefore, we will use the FileSource connector. We must create a configuration file to use with this connector. For this most part you can copy the example available in `$KAFKA_HOME/config/connect-file-source.properties`. Below is an example of our `my-file-source.properties` file.

Open a new file.

```
#vi my-file-source.properties
```

Paste the following instruction in the above file.

```
#my-file-source.properties config file
name=local-file-source
connector.class=com.github.mmolimar.kafka.connect.fs.FsSourceConnector
policy.class=com.github.mmolimar.kafka.connect.fs.policy.SimplePolicy
tasks.max=1
fs.uris=file:///tmp/my-test.txt
topic=my-kafka-connect
file_reader.class=com.github.mmolimar.kafka.connect.fs.file.reader.TextFileReader
file_reader.batch_size=0
```

This file indicates that we will use the FileStreamSource connector class, read data from the /tmp/my-test.txt file, and publish records to the my-kafka-connect Kafka topic. We are also only using 1 task to push this data to Kafka, since we are reading/publishing a single file.

Creating a Worker Config File.

Processes that execute Kafka Connect connectors and tasks are called workers. Since we are reading data from a single machine and publishing to Kafka, we can use the simpler of the two types, standalone workers (as opposed to distributed workers). You can find a sample config file for standalone workers in \$KAFKA_HOME/config/connect-standalone.properties. We will call our file my-standalone.properties.

Create a file.

```
#vi my-standalone.properties
```

Update with the following contents.

```
#bootstrap kafka servers  
bootstrap.servers=kafka:9092
```

```
# specify input data format  
key.converter=org.apache.kafka.connect.storage.StringConverter  
value.converter=org.apache.kafka.connect.storage.StringConverter
```

```
# local file storing offsets and config data  
offset.storage.file.filename=/tmp/connect.offsets  
plugin.path=/software/plugins
```

The main change in this example in comparison to the default is the key.converter and value.converter settings. Since our file contains simple text, we use the StringConverter types.

Download the file connector.

<https://www.confluent.io/hub/mmolimar/kafka-connect-fs>

Installation

Confluent Hub CLI installation

Use the [Confluent Hub client](#) to install this connector with:

```
$ confluent-hub install mmolimar/kafka-connect-fs:1.3.0
```

 [Copy](#)

Download installation

Or download the ZIP file and extract it into one of the directories that is listed on the Connect worker's plugin.path configuration properties. This must be done on each of the installations where Connect will be run.

[Download](#)

By downloading you agree to the [terms of use](#) and [software license agreement](#).

Download the above file.

```
#mkdir -p /software/plugins/
```

Rename the folder:

```
#mv mmolimar-kafka-connect-fs-1.3.0/ /software/plugins/kafka-file
```

Open a terminal.

Our input file /tmp/my-test.txt will be read in a single process to the Kafka my-kafka-connect topic. Here is a look at the file contents:

Create a file and paste the following text.

```
#vi /tmp/my-test.txt
```

This Message is from Test File.

It will be consumed by the Kafka Connector.

Running Kafka Connect.

Now it is time to run Kafka Connect with our worker and source configuration files. As mentioned before we will be running Kafka Connect in standalone mode. Here is an example of doing this with our custom config files:

```
#/opt/kafka/bin/connect-standalone.sh my-standalone.properties my-file-source.properties
```

```
mSourceTask, sleeping to wait for it to be created (org.apache.kafka.connect.fil  
e.FileStreamSourceTask:108)  
[2018-06-17 12:55:35,218] WARN Couldn't find file /tmp/my-test.txt for FileStreamSourceTask, sleeping to wait for it to be created (org.apache.kafka.connect.fil  
e.FileStreamSourceTask:108)  
[2018-06-17 12:55:36,226] WARN Couldn't find file /tmp/my-test.txt for FileStreamSourceTask, sleeping to wait for it to be created (org.apache.kafka.connect.fil  
e.FileStreamSourceTask:108)  
[2018-06-17 12:55:37,241] WARN Couldn't find file /tmp/my-test.txt for FileStreamSourceTask, sleeping to wait for it to be created (org.apache.kafka.connect.fil  
e.FileStreamSourceTask:108)  
[2018-06-17 12:55:38,250] WARN Couldn't find file /tmp/my-test.txt for FileStreamSourceTask, sleeping to wait for it to be created (org.apache.kafka.connect.fil  
e.FileStreamSourceTask:108)  
[2018-06-17 12:55:39,256] WARN Couldn't find file /tmp/my-test.txt for FileStreamSourceTask, sleeping to wait for it to be created (org.apache.kafka.connect.fil  
e.FileStreamSourceTask:108)  
[2018-06-17 12:55:40,258] WARN Couldn't find file /tmp/my-test.txt for FileStreamSourceTask, sleeping to wait for it to be created (org.apache.kafka.connect.fil  
e.FileStreamSourceTask:108)  
[2018-06-17 12:55:41,272] WARN Couldn't find file /tmp/my-test.txt for FileStreamSourceTask, sleeping to wait for it to be created (org.apache.kafka.connect.fil  
e.FileStreamSourceTask:108)
```

```
[2018-06-17 12:56:53,991] WARN Couldn't find file /tmp/my-test.txt for FileStreamSourceTask, sleeping to wait for it to be created (org.apache.kafka.connect.file.FileStreamSourceTask:108)
[2018-06-17 12:56:55,535] INFO Cluster ID: rwk6R3n9TYSCArCsSs8V4w (org.apache.kafka.clients.Metadata:265)
[2018-06-17 12:57:53,852] INFO WorkerSourceTask{id=local-file-source-0} Committing offsets (org.apache.kafka.connect.runtime.WorkerSourceTask:328)
[2018-06-17 12:57:53,854] INFO WorkerSourceTask{id=local-file-source-0} flushing 0 outstanding messages for offset commit (org.apache.kafka.connect.runtime.WorkerSourceTask:345)
[2018-06-17 12:57:54,062] INFO WorkerSourceTask{id=local-file-source-0} Finished commitOffsets successfully in 208 ms (org.apache.kafka.connect.runtime.WorkerSourceTask:427)
```

Open another terminal and execute the following consumer to consume the message.

Reading from the Kafka Topic

If we read from the Kafka topic that we created earlier, we should see the 2 lines in the source file that were written to Kafka:

```
#/opt/kafka/bin/kafka-console-consumer.sh --bootstrap-server kafka:9092 --topic my-kafka-connect --from-beginning
```

```
[root@tos ~]# /opt/kafka/bin/kafka-console-consumer.sh --bootstrap-server tos.master.com:9092 --topic my-kafka-connect --from-beginning
This Message is from Test File.
It will be consumed by the Kafka Connector.
```

We have successfully configured, file connector in standalone mode.

Next, let us configure JDBC in a distributed mode.

10. Kafka Connector (JDBC) - 60 Minutes

Kafka Connect – JDBC Connector.

Download the connector from the following url.

Tools - <https://www.confluent.io/hub/confluentinc/kafka-connect-jdbc>

Download installation

Or download the ZIP file and extract it into one of the directories that is listed on the Connect worker's plugin.path configuration properties. This must be done on each of the installations where Connect will be run.

[Download](#)

Extract the file:

```
# yum install unzip  
# unzip confluent*zip
```

```
[root@4f5607f478bd software]# pwd  
/software  
[root@4f5607f478bd software]# ls  
README.md  confluentinc-kafka-connect-jdbc-10.0.0_hbase-2.3.2-bin.tar  log.txt  master.txt  
[root@4f5607f478bd software]# ls confluentinc-kafka-connect-jdbc-10.0.0/  
assets doc etc lib manifest.json  
[root@4f5607f478bd software]#
```

Rename the folder:

```
#mv confluentinc-kafka-connect-jdbc-10.0.0/ /software/plugins/kafka-jdbc
```

Add this to the plugin path in your Connect properties file.

For example, `plugin.path=/software/plugins`. Kafka Connect finds the plugins using its plugin path.

Start the Connect workers with that configuration. Connect will discover all connectors defined within those plugins.

We can create /software/connect-distributed.properties file to specify the worker properties as follows:

Note that the `plugin.path` is the path that we need to place the library that we downloaded.

```
#vi /software/connect-distributed.properties
```

Update the following content in the above file.

```
# A list of host/port pairs to use for establishing the initial connection to the Kafka
cluster.
bootstrap.servers=kafka0:9092

# unique name for the cluster, used in forming the Connect cluster group. Note that this
must not conflict with consumer group IDs
group.id=connect-cluster

# The converters specify the format of data in Kafka and how to translate it into Connect
data.
key.converter=org.apache.kafka.connect.json.JsonConverter
value.converter=org.apache.kafka.connect.json.JsonConverter
key.converter.schemas.enable=true
value.converter.schemas.enable=true

# Topic to use for storing offsets. This topic should have many partitions and be replicated
and compacted.
offset.storage.topic=connect-offsets
offset.storage.replication.factor=1

# Topic to use for storing connector and task configurations; note that this should be a
single partition, highly replicated,
config.storage.topic=connect-configs
config.storage.replication.factor=1

# Topic to use for storing statuses. This topic can have multiple partitions and should be
replicated and compacted.
status.storage.topic=connect-status
status.storage.replication.factor=1

# Flush much faster than normal, which is useful for testing/debugging
offset.flush.interval.ms=10000

plugin.path=/software/plugins
```

Note that the `plugin.path` is the path that we need to place the library that we downloaded.

Start the kafka connector

Go to the bin folder of the kafka installation or specify the full path.

```
#/opt/kafka/bin/connect-distributed.sh /software/connect-distributed.properties
```

```
nnectorIds=[], taskIds=[], revokedConnectorIds=[], revokedTaskIds=[], delay=0} with rebalance delay: 0 (org.apache.kafka.connect.runtime.distr
ibuted.DistributedHerder:1681)
[2020-11-15 12:59:43,196] INFO [Worker clientId=connect-1, groupId=connect-cluster] Starting connectors and tasks using config offset -1 (org.
apache.kafka.connect.runtime.distributed.DistributedHerder:1208)
[2020-11-15 12:59:43,196] INFO [Worker clientId=connect-1, groupId=connect-cluster] Finished starting connectors and tasks (org.apache.kafka.c
onnect.runtime.distributed.DistributedHerder:1236)
Nov 15, 2020 12:59:43 PM org.glassfish.jersey.internal.Errors logErrors
WARNING: The following warnings have been detected: WARNING: The (sub)resource method listLoggers in org.apache.kafka.connect.runtime.res
ources.LoggingResource contains empty path annotation.
WARNING: The (sub)resource method createConnector in org.apache.kafka.connect.runtime.rest.resources.ConnectorsResource contains empty path an
notation.
WARNING: The (sub)resource method listConnectors in org.apache.kafka.connect.runtime.rest.resources.ConnectorsResource contains empty path ann
otation.
WARNING: The (sub)resource method listConnectorPlugins in org.apache.kafka.connect.runtime.rest.resources.ConnectorPluginsResource contains em
pty path annotation.
WARNING: The (sub)resource method serverInfo in org.apache.kafka.connect.runtime.rest.resources.RootResource contains empty path annotation.

[2020-11-15 12:59:43,307] INFO Started o.e.j.s.ServletContextHandler@c6da8bb{/null,AVAILABLE} (org.eclipse.jetty.server.handler.ContextHandle
r:825)
[2020-11-15 12:59:43,307] INFO REST resources initialized; server is started and ready to handle requests (org.apache.kafka.connect.runtime.re
st.RestServer:319)
[2020-11-15 12:59:43,308] INFO Kafka Connect started (org.apache.kafka.connect.runtime.Connect:57)
[2020-11-15 12:59:43,419] INFO [Worker clientId=connect-1, groupId=connect-cluster] Session key updated (org.apache.kafka.connect.runtime.dist
ributed.DistributedHerder:1570)
```

After running the connector we can confirm that connector's REST endpoint is accessible, and we can confirm that JDBC connector is in the plugin list by calling <http://localhost:8083/connector-plugins>

```
# curl http://localhost:8083/connector-plugins
```

```
[root@4f5607f478bd software]# curl http://localhost:8083/connector-plugins
[{"class":"io.confluent.connect.jdbc.JdbcSinkConnector","type":"sink","version":"10.0.0"}, {"class":"io.confluent.connect.jdbc.JdbcSourceConnector","type":"source","version":"10.0.0"}, {"class":"org.apache.kafka.connect.file.FileStreamSinkConnector","type":"sink","version":"2.6.0"}, {"class":"org.apache.kafka.connect.file.FileStreamSourceConnector","type":"source","version":"2.6.0"}, {"class":"org.apache.kafka.connect.mirror.MirrorCheckpointConnector","type":"source","version":"1"}, {"class":"org.apache.kafka.connect.mirror.MirrorHeartbeatConnector","type":"source","version":"1"}, {"class":"org.apache.kafka.connect.mirror.MirrorSourceConnector","type":"source","version":"1"}][root@4f5607f478bd software]#
[root@4f5607f478bd software]#
```

Install the postgres db. We will configure postgresdb to store some data which will be transfer by connect to kafka topic.

Refer url for downloading the postgresdb software : <https://www.postgresql.org/download/>.

For Centos 7:

[If you are using docker, initiate a container for PostgresDB:

```
#docker run -it --name postgresdb --privileged -p 5432:5432 centos:7 /usr/sbin/init
```

```
# docker exec -it postgresdb bash
```

Add both the container in same network.

```
#docker network connect spark-net postgresdb
```

```
]
```

The following command will install and initialize the postgres DB

```
#yum install -y https://download.postgresql.org/pub/repos/yum/reporpms/EL-7-x86\_64/pgdg-redhat-repo-latest.noarch.rpm
```

Or

```
# yum install -y https://download.postgresql.org/pub/repos/yum/reporpms/EL-8-aarch64/pgdg-redhat-repo-latest.noarch.rpm

#yum install -y postgresql13-server
#/usr/pgsql-13/bin/postgresql-13-setup initdb
```

Enable and start the service:

```
systemctl enable postgresql-13
systemctl start postgresql-13
```

```
Complete!
[root@5a1225a251e4 ~]# /usr/pgsql-13/bin/postgresql-13-setup initdb
Initializing database ... OK

[root@5a1225a251e4 ~]# systemctl enable postgresql-13
Created symlink from /etc/systemd/system/multi-user.target.wants/postgresql-13.service to /usr/lib/systemd/system/postgresql-13.service.
[root@5a1225a251e4 ~]# systemctl start postgresql-13
[root@5a1225a251e4 ~]# systemctl status postgresql-13
● postgresql-13.service - PostgreSQL 13 database server
  Loaded: loaded (/usr/lib/systemd/system/postgresql-13.service; enabled; vendor preset: disabled)
  Active: active (running) since Sun 2020-11-15 13:21:23 UTC; 6s ago
    Docs: https://www.postgresql.org/docs/13/static/
   Process: 322 ExecStartPre=/usr/pgsql-13/bin/postgresql-13-check-db-dir ${PGDATA} (code=exited, status=0/SUCCESS)
 Main PID: 327 (postmaster)
  CGroup: /docker/5a1225a251e4011bfe46ec657ce6ede8168c9c54e77cecc29df0c095aeebbd37/docker/5a1225a251e4011bfe46ec657ce6ede8168c9c54e77cecc29df0c095aeebbd37/system.slice/postgresql-13.service
          ├ 327 /usr/pgsql-13/bin/postmaster -D /var/lib/pgsql/13/data/
          ├ 328 postgres: logger
          ├ 330 postgres: checkpointer
          └ 331 postares: background writer
```


The configuration for the plugin is stored in /software/jdbc-source.json file on kafka broker server. It contents is as follows: (replace localhost with container IP of postgress DB if you are using docker)

```
{  
  "name": "jdbc_source_connector_postgresql_01",  
  "config": {  
    "connector.class": "io.confluent.connect.jdbc.JdbcSourceConnector",  
    "connection.url": "jdbc:postgresql://localhost:5432/postgres",  
    "connection.user": "postgres",  
    "connection.password": "postgres",  
    "topic.prefix": "postgres-01-",  
    "poll.interval.ms": 3600000,  
    "mode": "bulk"  
  }  
}
```

Start Postgress DB CLI: On the postgres DB server.

```
sudo -u postgres psql  
ALTER USER postgres PASSWORD 'postgres';
```

```
[root@5a1225a251e4 ~]# sudo -u postgres psql  
psql (13.1)  
Type "help" for help.  
  
postgres=# ALTER USER postgres PASSWORD 'postgres';  
ALTER ROLE  
postgres=#
```

if sudo doesn't exists in your system, install (yum install sudo)

Create a table and insert few records: Execute all the bellows SQL commands from the Postgresdb CLI.

```
#CREATE TABLE leads (id INTEGER PRIMARY KEY, name VARCHAR);
insert into leads values (1,'Henry P');
insert into leads values (2,'Rajnita P');
insert into leads values (3,'Henderson P');
insert into leads values (4,'Tiraj P');
select * from leads;
```

```
postgres=# insert into leads values (4,'Tiraj P');
INSERT 0 1
postgres=# select * from leads;
 id |      name
----+
  1 | Henry P
  2 | Rajnita P
  3 | Henderson P
  4 | Tiraj P
(4 rows)
```

To exit : \q

Allow all to connect to postgresDB from remote server by updating the following contents:

```
#vi /var/lib/pgsql/13/data/postgresql.conf
```

```
listen_addresses = '*'
```

Append in the last line of the following file.

```
# vi /var/lib/pgsql/13/data/pg_hba.conf
```

```
host all all all trust
```

```
#systemctl restart postgresql-13
```

Starting the JDBC Connector – On the Kafka Node.

As we operate on distributed mode we run the connectors by calling REST endpoints with the configuration JSON. We can specify the configuration payload from a file for curl command. The following command starts the connector. Execute the following command from the directory you have stored the configuration json file.

```
curl -d @"jdbc-source.json" \  
-H "Content-Type: application/json" \  
-X POST http://localhost:8083/connectors
```

```
timestamp.column.name = □
timestamp.delay.interval.ms = 0
timestamp.initial = null
topic.prefix = postgres-02-
validate.non.null = true
(io.confluent.connect.jdbc.source.JdbcSourceTaskConfig:354)
[2020-11-15 16:05:29,313] INFO Using JDBC dialect PostgreSQL (io.confluent.connect.jdbc.source.JdbcSourceTask:98)
[2020-11-15 16:05:29,317] INFO Attempting to open connection #1 to PostgreSQL (io.confluent.connect.jdbc.util.CachedConnectionProvider:82)
[2020-11-15 16:05:29,383] INFO Started JDBC source task (io.confluent.connect.jdbc.source.JdbcSourceTask:257)
[2020-11-15 16:05:29,384] INFO WorkerSourceTask{id=jdbc_source_connector_postgresql_04-0} Source task finished initialization and start (org.apache.kafka.connect.runtime.WorkerSourceTask:233)
[2020-11-15 16:05:29,388] INFO Begin using SQL query: SELECT * FROM "public"."leads" (io.confluent.connect.jdbc.source.TableQuerier:164)
[2020-11-15 16:05:29,503] WARN [Producer clientId=connector-producer-jdbc_source_connector_postgresql_04-0] Error while fetching metadata with correlation id 3 : {postgres-02-leads=LEADER_NOT_AVAILABLE} (org.apache.kafka.clients.NetworkClient:1073)
[2020-11-15 16:05:39,266] INFO WorkerSourceTask{id=jdbc_source_connector_postgresql_04-0} Committing offsets (org.apache.kafka.connect.runtime.WorkerSourceTask:478)
[2020-11-15 16:05:39,268] INFO WorkerSourceTask{id=jdbc_source_connector_postgresql_04-0} flushing 0 outstanding messages for offset commit (org.apache.kafka.connect.runtime.WorkerSourceTask:495)
[2020-11-15 16:05:39,295] INFO WorkerSourceTask{id=jdbc_source_connector_postgresql_04-0} Finished commitOffsets successfully in 28 ms (org.apache.kafka.connect.runtime.WorkerSourceTask:574)
[2020-11-15 16:05:49,264] INFO WorkerSourceTask{id=jdbc_source_connector_postgresql_04-0} Committing offsets (org.apache.kafka.connect.runtime.WorkerSourceTask:478)
[2020-11-15 16:05:49,265] INFO WorkerSourceTask{id=jdbc_source_connector_postgresql_04-0} flushing 0 outstanding messages for offset commit (org.apache.kafka.connect.runtime.WorkerSourceTask:495)
```

We can see that in postgres database table leads is loaded to kafka topic- postgres-02-leads:

```
#/opt/kafka/bin/kafka-topics.sh --list --bootstrap-server kafka0:9092
```

```
#/opt/kafka/bin/kafka-topics.sh --describe --bootstrap-server kafka0:9092 --topic postgres-01-leads
```

```
Topic: connect-status Partition: 4 Leader: 0 Replicas: 0 Isr: 0
in/ Topic: postgres-02-leads PartitionCount: 1 ReplicationFactor: 1 Configs:
in/ Topic: postgres-02-leads Partition: 0 Leader: 0 Replicas: 0 Isr: 0
in/ Topic: test PartitionCount: 1 ReplicationFactor: 1 Configs:
in/ Topic: test Partition: 0 Leader: 0 Replicas: 0 Isr: 0
[root@4f5607f478bd software]#
```

And each row in the tables are loaded as a message. You can verify using the consumer console.

```
#/opt/kafka/bin/kafka-console-consumer.sh --bootstrap-server kafka0:9092 --topic postgres-01-leads --from-beginning
```

```
[root@4f5607f478bd software]# /opt/kafka/bin/kafka-console-consumer.sh --bootstrap-server localhost:9092 --topic postgres-02-leads --from-beginning
{"schema": {"type": "struct", "fields": [{"type": "int32", "optional": false, "field": "id"}, {"type": "string", "optional": true, "field": "name"}]}, "optional": false, "name": "leads"}, "payload": {"id": 1, "name": "Henry P"}}
{"schema": {"type": "struct", "fields": [{"type": "int32", "optional": false, "field": "id"}, {"type": "string", "optional": true, "field": "name"}]}, "optional": false, "name": "leads"}, "payload": {"id": 2, "name": "Rajnita P"}}
{"schema": {"type": "struct", "fields": [{"type": "int32", "optional": false, "field": "id"}, {"type": "string", "optional": true, "field": "name"}]}, "optional": false, "name": "leads"}, "payload": {"id": 3, "name": "Henderson P"}}
{"schema": {"type": "struct", "fields": [{"type": "int32", "optional": false, "field": "id"}, {"type": "string", "optional": true, "field": "name"}]}, "optional": false, "name": "leads"}, "payload": {"id": 4, "name": "Tiraj P"}}
```

Let us perform a Single Message Transformation.

Scenario: **Name** field should be changed to **FullName**

Solution: We will use the ReplaceField to perform the above activity.

Add another instance of connector.

Define the Single Message Transformation parameters.

The configuration for the plugin is stored in /software/jdbc-source_smt.json file on kafka broker server. It contents is as follows: (replace localhost with container IP of postgress DB if you are using docker)

```
{  
  "name": "jdbc_source_connector_postgresql_02",  
  "config": {  
    "connector.class": "io.confluent.connect.jdbc.JdbcSourceConnector",  
    "connection.url": "jdbc:postgresql://localhost:5432/postgres",  
    "connection.user": "postgres",  
    "connection.password": "postgres",  
    "topic.prefix": "postgres-02-",  
    "poll.interval.ms": 3600000,  
    "mode": "bulk",  
    "transforms": "RenameField",  
    "transforms.RenameField.type": "org.apache.kafka.connect.transforms.ReplaceField$Value",  
    "transforms.RenameField.renames": "name:fullname"  
  }  
}
```

In the above configuration, the field , name is replace with fullname attribute.

Start the instance

```
#curl -d @"jdbc-source_smt1.json" \  
-H "Content-Type: application/json" \  
-X POST http://localhost:8083/connectors
```

Verify the topic and check the message inside that topic.

```
#/opt/kafka/bin/kafka-topics.sh --list --bootstrap-server kafka0:9092  
#/opt/kafka/bin/kafka-console-consumer.sh --bootstrap-server kafka0:9092 --topic postgres-02-leads --from-beginning
```

```
[root@kafka0 software]# /opt/kafka/bin/kafka-topics.sh --list --bootstrap-server localhost:9092  
__consumer_offsets  
connect-configs  
connect-offsets  
connect-status  
postgres-01-leads  
postgres-02-leads  
[root@kafka0 software]# /opt/kafka/bin/kafka-console-consumer.sh --bootstrap-server localhost:9092 --topic postgres-02-leads --from-beginning  
{"schema": {"type": "struct", "fields": [{"type": "int32", "optional": false, "field": "id"}, {"type": "string", "optional": true, "field": "fullname"}], "optional": false, "name": "leads"}, "payload": {"id": 1, "fullname": "Henry P"}}  
{"schema": {"type": "struct", "fields": [{"type": "int32", "optional": false, "field": "id"}, {"type": "string", "optional": true, "field": "fullname"}], "optional": false, "name": "leads"}, "payload": {"id": 2, "fullname": "Rajnita P"}}  
{"schema": {"type": "struct", "fields": [{"type": "int32", "optional": false, "field": "id"}, {"type": "string", "optional": true, "field": "fullname"}], "optional": false, "name": "leads"}, "payload": {"id": 3, "fullname": "Henderson P"}}  
{"schema": {"type": "struct", "fields": [{"type": "int32", "optional": false, "field": "id"}, {"type": "string", "optional": true, "field": "fullname"}], "optional": false, "name": "leads"}, "payload": {"id": 4, "fullname": "Tiraj P"}}
```

As shown above the message field has been transform from Name to fullname.

Execute some Queries to get information about connectors:

Get a list of active connectors.

```
#curl http://localhost:8083/connectors
```

Get configuration info for a connector.

```
#curl http://localhost:8083/connectors/jdbc\_source\_connector\_postgresql\_01/config
```

```
[root@kafka0 /]# curl http://localhost:8083/connectors
[{"connector.class": "io.confluent.connect.jdbc.JdbcSourceConnector", "mode": "bulk", "topic.prefix": "postgres-01-", "connection.password": "postgres", "connection.user": "postgres", "poll.interval.ms": "3600000", "name": "jdbc_source_connector_postgresql_01", "connection.url": "jdbc:postgresql://localhost:5432/postgres"}][root@kafka0 /]#
[root@kafka0 /]# curl http://localhost:8083/connectors/jdbc_source_connector_postgresql_01/config
{"connector.class": "io.confluent.connect.jdbc.JdbcSourceConnector", "mode": "bulk", "topic.prefix": "postgres-01-", "connection.password": "postgres", "connection.user": "postgres", "poll.interval.ms": "3600000", "name": "jdbc_source_connector_postgresql_01", "connection.url": "jdbc:postgresql://localhost:5432/postgres"}[root@kafka0 /]#
[root@kafka0 /]# curl http://localhost:8083/connectors/jdbc_source_connector_postgresql_02/config
{"connector.class": "io.confluent.connect.jdbc.JdbcSourceConnector", "mode": "bulk", "transforms.RenameField.renames": "name:fullname", "topic.prefix": "postgres-02-", "connection.password": "postgres", "connection.user": "postgres", "poll.interval.ms": "3600000", "transforms.RenameField.type": "org.apache.kafka.connect.transforms.ReplaceField$Value", "transforms": "RenameField", "name": "jdbc_source_connector_postgresql_02", "connection.url": "jdbc:postgresql://localhost:5432/postgres"}[root@kafka0 /]#
[root@kafka0 /]# curl http://localhost:8083/connectors/jdbc_source_connector_postgresql_02/tasks
[{"id": {"connector": "jdbc_source_connector_postgresql_02", "task": 0}, "config": {"connector.class": "io.confluent.connect.jdbc.JdbcSourceConnector", "transforms.RenameField.renames": "name:fullname", "connection.password": "postgres", "transforms": "RenameField", "mode": "bulk", "topic.prefix": "postgres-02-", "tables": "\\"public\\\".\\"leads\\\"", "task.class": "io.confluent.connect.jdbc.source.JdbcSourceTask", "connection.user": "postgres", "poll.interval.ms": "3600000", "transforms.RenameField.type": "org.apache.kafka.connect.transforms.ReplaceField$Value", "name": "jdbc_source_connector_postgresql_02", "connection.url": "jdbc:postgresql://localhost:5432/postgres"}][root@kafka0 /]#
```

Retrieve details for specific tasks (JDBC Tasks) - Get a list of tasks currently running for the connector.

```
#curl http://localhost:8083/connectors/jdbc\_source\_connector\_postgresql\_02/tasks
```

```
[root@kafka0 ~]# curl http://localhost:8083/connectors/jdbc_source_connector_postgresql_02/tasks
[{"id":{"connector":"jdbc_source_connector_postgresql_02","task":0}, "config":{"connector.class":"io.confluent.connect.jdbc.JdbcSourceConnector", "transforms.RenameField.renames":"name:fullname", "connection.password":"postgres", "transforms":"Rename Field", "mode":"bulk", "topic.prefix":"postgres-02-", "tables":"\"public\".\"leads\"", "task.class":"io.confluent.connect.jdbc.source.JdbcSourceTask", "connection.user":"postgres", "poll.interval.ms":"3600000", "transforms.RenameField.type":"org.apache.kafka.connect.transforms.ReplaceField$Value", "name":"jdbc_source_connector_postgresql_02", "connection.url":"jdbc:postgresql://localhost:5432/postgres"}]}][root@kafka0 ~]# curl http://localhost:8083/connectors/jdbc_source_connector_postgresql_01/tasks
[{"id":{"connector":"jdbc_source_connector_postgresql_01", "task":0}, "config":{"connector.class":"io.confluent.connect.jdbc.JdbcSourceConnector", "mode":"bulk", "topic.prefix":"postgres-01-", "tables":"\"public\".\"leads\"", "connection.password":"postgres", "task.class":"io.confluent.connect.jdbc.source.JdbcSourceTask", "connection.user":"postgres", "poll.interval.ms":"3600000", "name":"jdbc_source_connector_postgresql_01", "connection.url":"jdbc:postgresql://localhost:5432/postgres"}}][root@kafka0 ~]#
```

The current status of the connector

```
#curl http://localhost:8083/connectors/?expand=status
```

```
[root@kafka0 ~]# curl http://localhost:8083/connectors/?expand=status
{"jdbc_source_connector_postgresql_01": {"status": {"name": "jdbc_source_connector_postgresql_01", "connector": {"state": "RUNNING", "worker_id": "172.18.0.2:8083"}, "tasks": [{"id": 0, "state": "RUNNING", "worker_id": "172.18.0.2:8083"}], "type": "source"}}, "jdbc_source_connector_postgresql_02": {"status": {"name": "jdbc_source_connector_postgresql_02", "connector": {"state": "RUNNING", "worker_id": "172.18.0.2:8083"}, "tasks": [{"id": 0, "state": "RUNNING", "worker_id": "172.18.0.2:8083"}]}, "type": "source"}]}[root@kafka0 ~]#
```

Gets the current status of the connector, including:

- Whether it is running or restarting, or if it has failed or paused
- Which worker it is assigned to
- Error information if it has failed
- The state of all its tasks

```
#curl http://localhost:8083/connectors/jdbc\_source\_connector\_postgresql\_01/status
```

```
[root@kafka0 ~]# curl http://localhost:8083/connectors/jdbc_source_connector_postgresql_01/status
{"name":"jdbc_source_connector_postgresql_01","connector":{"state":"RUNNING","worker_id":"172.18.0.2:8083"},"tasks":[{"id":0,"state":"RUNNING","worker_id":"172.18.0.2:8083"}],"type":"source"}[root@kafka0 ~]#
```

Retrieve details for specific tasks

```
# curl
http://localhost:8083/connectors/jdbc\_source\_connector\_postgresql\_01/tasks/0/status
```

```
[root@kafka0 ~]# curl http://localhost:8083/connectors/jdbc_source_connector_postgresql_01/tasks/0/status
{"id":0,"state":"RUNNING","worker_id":"172.18.0.2:8083"}[root@kafka0 ~]#
```

----- Labs End Here -----

12. Errors

I. {test=LEADER_NOT_AVAILABLE} (org.apache.kafka.clients.NetworkClient)

```
[2018-05-15 23:46:40,132] WARN [Producer clientId=console-producer] Error while
fetching metadata with correlation id 14 : {test=LEADER_NOT_AVAILABLE} (org.apac
he.kafka.clients.NetworkClient)
[2018-05-15 23:46:40,266] WARN [Producer clientId=console-producer] Error while
fetching metadata with correlation id 15 : {test=LEADER_NOT_AVAILABLE} (org.apac
he.kafka.clients.NetworkClient)
^C[2018-05-15 23:46:40,394] WARN [Producer clientId=console-producer] Error whil
e fetching metadata with correlation id 16 : {test=LEADER_NOT_AVAILABLE} (org.apa
che.kafka.clients.NetworkClient)
[root@tos opt]# {test=LEADER_NOT_AVAILABLE} (org.apache.kafka.clients.NetworkCl
ient)
bash: syntax error near unexpected token `org.apache.kafka.clients.NetworkClient
'
```

Solutions: /opt/kafka/config/server.properties

Update the following information.

```
# it uses the value for "listeners" if configured. Otherwise, it will use the v
alue
# returned from java.net.InetAddress.getCanonicalHostName().
advertised.listeners=PLAINTEXT://localhost:9092
# Maps listener names to security protocols, the default is for them to be the s
ame as the listener name
```

II Unable to connect to the server / Topic my-example-topic not present in
metadata after 60000 ms.

```
##### Socket Server Settings #####
# The address the socket server listens on. It will get the value returned from
# java.net.InetAddress.getCanonicalHostName() if not configured.
#   FORMAT:
#     listeners = listener_name://host_name:port
#   EXAMPLE:
#     listeners = PLAINTEXT://your.host.name:9092
#listeners=PLAINTEXT://:9092

# Hostname and port the broker will advertise to producers and consumers. If not set,
# it uses the value for "listeners" if configured. Otherwise, it will use the value
# returned from java.net.InetAddress.getCanonicalHostName().
advertised.listeners=PLAINTEXT://localhost:9092

# Maps listener names to security protocols, the default is for them to be the same. See the config documentation for more details
#listener.security.protocol.map=PLAINTEXT:PLAINTEXT,SSL:SSL,SASL_PLAINTEXT:SASL_PLAINTEXT,SASL_SSL:SASL_SSL

"server.properties" 136L, 6848C written
```

13. LOG verification + segment Sizing

14. Annexure Code:

II. DumpLogSegment

```
/opt/kafka/bin/kafka-run-class.sh kafka.tools.DumpLogSegments --deep-iteration --print-data-log --files \
/tmp/kafka-logs/my-kafka-connect-0/oooooooooooooooooooo.log | head -n 4
```

```
[root@tos test-topic-0]# more 00000000000000000000.log
[root@tos test-topic-0]# cd ..
[root@tos kafka-logs]# cd my-kafka-connect-0/
[root@tos my-kafka-connect-0]# ls
00000000000000000000.index      00000000000000000011.snapshot
00000000000000000000.log        leader-epoch-checkpoint
00000000000000000000.timeindex
[root@tos my-kafka-connect-0]# more *log
\Kafka Connector.--More-- (53%)

[root@tos my-kafka-connect-0]# pwd
/tmp/kafka-logs/my-kafka-connect-0
[root@tos my-kafka-connect-0]# /opt/kafka/bin/kafka-run-class.sh kafka.tools.DumpLogSegments --deep-iteration --print-data-log --files \
> /tmp/kafka-logs/my-kafka-connect-0/00000000000000000000.log | head -n 4
Dumping /tmp/kafka-logs/my-kafka-connect-0/00000000000000000000.log
Starting offset: 0
offset: 0 position: 0 CreateTime: 1530552634675 isvalid: true keysize: -1 valuesize: 31 magic: 2 compresscodec: NONE producerId: -1 producerEpoch: -1 sequence: -1 isTransactional: false headerKeys: [] payload: This Message is from Test File .
offset: 1 position: 0 CreateTime: 1530552634677 isvalid: true keysize: -1 valuesize: 43 magic: 2 compresscodec: NONE producerId: -1 producerEpoch: -1 sequence: -1 isTransactional: false headerKeys: [] payload: It will be consumed by the Kafka Connector.
[root@tos my-kafka-connect-0]#
```

CLI to list offset number.

```
#/opt/kafka/bin/kafka-run-class.sh kafka.tools.GetOffsetShell --broker-list localhost:9092  
--topic IBM
```

Sending Message with Key and Value

```
#/opt/kafka/bin/kafka-console-producer.sh --broker-list localhost:9092 --topic IBM --  
property "parse.key=true" --property "key.separator=:"
```

```
#/opt/kafka/bin/Kafka-console-consumer.sh --topic IBM --bootstrap-server  
localhost:9092 --from-beginning \  
--property print.key=true \  
--property key.separator=":" \  
--partition 4 \  
--offset 3
```

```
kafka-console-consumer --topic example-topic --bootstrap-server broker:9092 \  
--from-beginning \  
--property print.key=true \  
--property key.separator="-" \  
--partition 0
```

III. Resources

<https://developer.ibm.com/hadoop/2017/04/10/kafka-security-mechanism-saslplain/>

<https://sharebigdata.wordpress.com/2018/01/21/implementing-sasl-plain/>

<https://developer.ibm.com/code/howtos/kafka-authn-authz>

<https://jaceklaskowski.gitbooks.io/apache-kafka/content/kafka-tools-GetOffsetShell.html>